



The Relational Model and the Structured Query Language (SQL)

Abdu Alawini

University of Illinois at Urbana-Champaign

CS411: Database Systems

June 12, 2020



Learning Objectives

After this lecture, you should be able to:

- Define a **data model**
- Distinguish among different **data models**
- Define the **relational data model**
- Articulate the **basic terminologies** of the **relational data model** (from a practical prospective)
- Define **Primary and Foreign keys** for a relational table
- Distinguish between **Schema** and **Instance**
- Write **single-relation SQL** queries



What is a Data Model?

A data model is a notation for **describing data or information**. The description generally consists of three parts:

1. Structure of the data:

- data structures used to implement data in the computer (physical data model)

2. Operations on the data:


- limited set of queries (operations that retrieve information) and modifications (operations that change the database).

3. Constraints on the data:

- ways to describe limitations on what the data can be. These constraints often come from the real-world application requirements


Types of Data Models

Hierarchical Model



Src: Wikipedia

Network Model



Src: Wikipedia

Semi-structured Model (XML)

```
<note>
  <date>2017-11-08</date>
  <hour>08:30</hour>
  <to>Raj</to>
  <from>Ravi</from>
  <body>Meeting at 8am.</body>
</note>
```



Outline

- ✓ Data Models
 - Relational Database Model
 - Basic Concepts and Terminology
 - Keys and Foreign Keys
 - Schema Specifications
 - SQL query language
 - Single-Relation Queries


The Relational Data Model

It all began with a breakthrough paper by E.F. Codd in 1970:
“A relational model of data for large shared data banks”.

Communications of the ACM 13 (6): 377

Codd's insights:

- Separate physical implementation from logical
- Model the data **independently** from how it will be used (accessed, printed, etc.)
 - Describe the data **minimally** and **mathematically**
 - A relation describes an association between data items – **tuples** with **attributes**
 - We generally think of tables and rows, but that's somewhat imprecise
- Use standard mathematical (logical) operations over the data – these are the **relational algebra** or **relational calculus**





Database from a user Perspective

- We'll assume that a DB has been already been implemented and loaded with data
- Our roles is to query/modify the data using SQL
- But before that, we need to learn the basics of relational model (from a practical point of view)



Introduction to Relational Databases from a Practical Point of View

Account

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

Imagine that this table (or relation) has been defined to help keep track of bank accounts.



Table Structure

The *name* of the table

Account

The name of the *columns* (*attributes*)

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking



Table Schema

The *schema* for the table

Account			
Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

The **schema** sets the structure of the table. You can think of the schema as the **definition** of the table. (Note, the schema specifies more information than what is shown.)



Table Rows

Account

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
...	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

Each entry in the table is called a **row** (**tuple**).
Sometimes an entry in the table is called a record.



Table Instance

An *instance* of the table...

the current contents or data in the table.

Account

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking



Another Table Instance

Another *instance* of the table
(two rows added, one (103) deleted)

Account

Number	Owner	Balance	Type
101	J. Smith	1,000.00	checking
102	W. Wei	2,000.00	checking
104	M. Jones	1,000.00	checking
105	H. Martin	10,000.00	checking
107	W. Yu	7,500.00	savings
109	R. Jones	432.55	checking

new



Intension vs. Extension

The *intension* of the table

Account

Number	Owner	Balance	Type
101	J. Smith	1000.00	checking
102	W. Wei	2000.00	checking
103	J. Smith	5000.00	savings
104	M. Jones	1000.00	checking
105	H. Martin	10,000.00	checking

The *extension* of the table. Also called the *extent*.



“Size” of a Table

Degree or arity of a table is the number of columns

Degree of this relation (or table) is 4
because there are 4 attributes

Account	Number	Owner	Balance	Type
→	101	J. Smith	1000.00	checking
→	102	W. Wei	2000.00	checking
→	103	J. Smith	5000.00	savings
→	104	M. Jones	1000.00	checking
→	105	H. Martin	10,000.00	checking

Cardinality
of this instance
is 5 (because
there are 5
rows)

Cardinality of a table = the number of rows in the
current instance



Outline

- ✓ Data Models
- Relational Database Model
 - ✓ Basic Concepts and Terminology
 - Keys and Foreign Keys
 - Schema Specifications
- SQL query language
 - Single-Relation Queries



Database (One or More Tables)

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00

Check	AcctNo	Check-number	Date	Amount
	101	924	10/23/00	125.00
	101	925	10/24/00	23.98



Table Keys

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00

Check	AcctNo	Check-number	Date	Amount
	101	924	10/23/00	125.00
	101	925	10/24/00	23.98

Each table has a key.... where the values must be unique.



Table Keys (cont.)

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00

Check	AcctNo	Check-number	Date	Amount
	101	924	10/23/00	125.00
	101	925	10/24/00	23.98

Key may consist of one column or two (or more) columns.



Connections between Tables

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00
	106	5	12/05/00	555.00

Is this legal?

If not, how do we prevent it from happening?



Foreign Key

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking

Deposit	AcctNo	Transaction-id	Date	Amount
	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00
	106	5	12/05/00	555.00

We say that **Deposit.AcctNo** is a *foreign key* that *references* **Account.Number**. If the DBMS enforces this constraint, we have *referential integrity*.



Foreign Keys (cont.)

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking
Check	AcctNo	Check-number	Date	Amount
	101	924	10/23/98	125.00
	101	925	10/24/98	23.98

Are there any foreign keys in the Check table?

Yes, Check.AcctNo is a foreign key that references Account.Number.



Foreign keys might or might not be part of the key for the referring table

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking

Deposit	AcctNo	Transaction-id	Date	Amount
Deposit.AcctNo is not part of key for Deposit.	102	1	10/22/00	500.00
	102	2	10/29/00	200.00
	104	3	10/29/00	1000.00
	105	4	11/02/00	10,000.00

Check	AcctNo	Check-number	Date	Amount
Check.AcctNo is part of key for Check.	101	924	10/23/00	125.00
	101	925	10/24/00	23.98

Foreign Key

Primary Key



Outline

- ✓ Data Models
- Relational Database Model
 - ✓ Basic Concepts and Terminology
 - ✓ Keys and Foreign Keys
 - Schema Specifications
- SQL query language
 - Single-Relation Queries



Specification of a Database Schema

- Select the tables, with a **name for each table**.
- Select **columns for each table** and give the **domain for each column**.
- Specify the **key(s)** for each table.

There can be more than one key for a table.
- Specify all appropriate **foreign keys**.



Database Domains for Columns

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking

...

For every column of every table, **the schema specifies allowable values**. For example,

Number must be a 3-digit number

Owner must be a 30-character string

Type must be “checking” or “savings”

The set of allowable values for a column is called the **domain** of the column.



Example Database Schema

(Keys are underlined. Each table has one key.)

Student		Takes				Course		
<u>sid</u>	name	<u>sid</u>	exp-grade	<u>cid</u>	<u>sem</u>	<u>cid</u>	subj	<u>sem</u>

Professor		Teaches		
<u>fid</u>	name	<u>fid</u>	<u>cid</u>	<u>sem</u>

In relational DBs, we use *relation(attribute:domain)*

STUDENT(sid:int, name:string)

Takes(sid:int, exp-grade:char[2], cid:string, sem:char[3])

COURSE(cid:string, subj:string, sem:char[3])

Teaches(fid:int, cid:string, sem:char[3])

PROFESSOR(fid:int, name:string)



Outline

- ✓ Data Models
- Relational Database Model
 - ✓ Basic Concepts and Terminology
 - ✓ Keys and Foreign Keys
 - ✓ Schema Specifications
- SQL query language
 - Single-Relation Queries

SQL: Structured Query Language



The standard language for relational data

- Invented by folks at IBM, esp. Don Chamberlin
- Actually not a particularly elegant language...
- Beat a more elegant competing standard, QUEL, from Berkeley

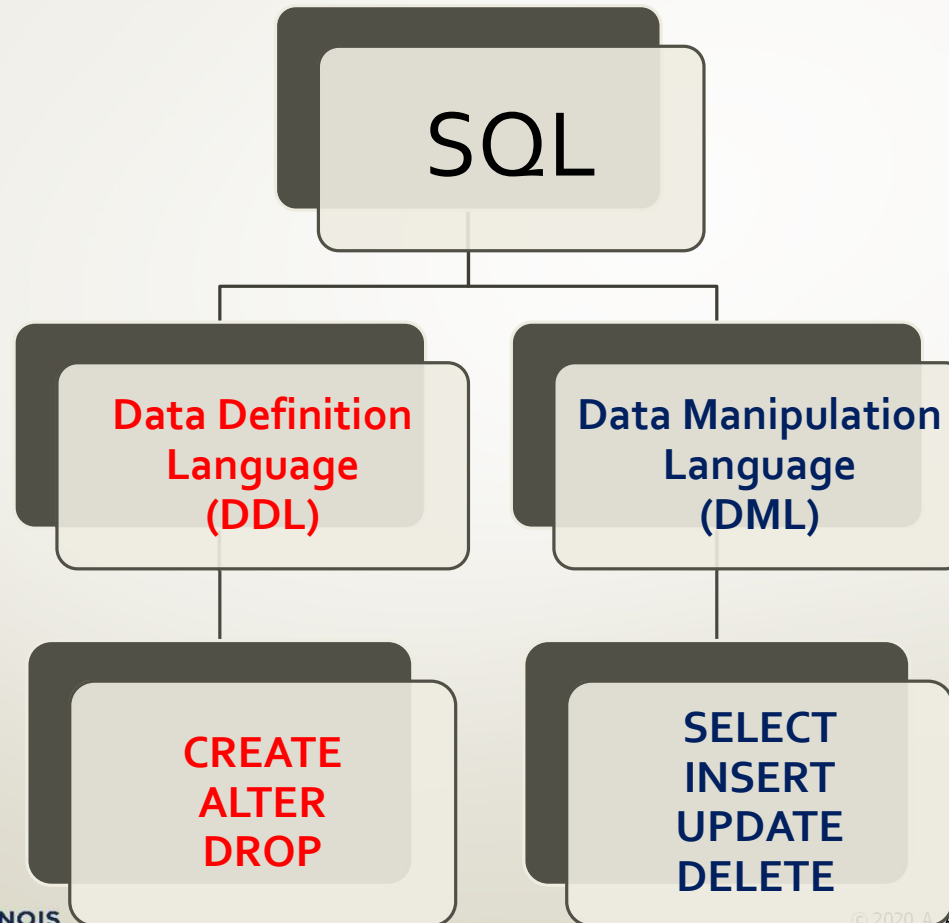
Separated into

- DDL (data definition language)
- DML (data manipulation language)

DML based on **relational calculus**, which we discuss later



SQL – the language we use to talk to the Database Management System





SQL (cont.)

SQL is a standard...
and there have been a series of SQL standards:
1986, 1989, 1992 (SQL2), 1999 (SQL3), ...,
SQL:2011

But DBMS products differ in how much of the
standard they support ... and how many extra
features they have.



Single-Relation Queries

Structured Query Language



Single-Relation Query Form

- The principal form of a single-relation query is:

SELECT desired attributes

FROM one table (relation)

WHERE condition about tuples of the table



Example Query

Account	Number	Owner	Balance	Type
	101	J. Smith	1000.00	checking
	102	W. Wei	2000.00	checking
	103	J. Smith	5000.00	savings
	104	M. Jones	1000.00	checking
	105	H. Martin	10,000.00	checking

```
SELECT      Number, Owner
FROM        Account
WHERE       Type = "savings";
```

Number	Owner
103	J. Smith



How a Single-Relation SQL query is evaluated

Third, the SELECT clause tells us which columns to keep in the query answer.

SELECT
FROM
WHERE

Number, Owner

Account

Type = "savings"

First, the FROM clause tells us the input tables.

Second, the WHERE clause is evaluated for all rows from the input table.



Renaming Attributes

- If you want the result to have different attribute names, use “AS <new name>” to rename an attribute.
- Example based on Account(number, owner, balance, type):

```
SELECT Number AS Acc_Num, Owner  
FROM Accounts  
WHERE Type = "savings";
```

Acc_Num	Owner
103	J. Smith



Another Database Schema

- We will be using the following database schema as a running example.

- Underline indicates key attributes.

Drinks(name, manf)

Cafe(name, addr, license)

Customers(name, addr, phone)

Likes(customer, drink)

Sells(cafe, drink, price)

Frequents(customer, cafe)



Expressions in SELECT Clauses

- Any expression that makes sense can appear as an element of a SELECT clause.
- Example: from Sells (cafe, drink, price):

```
SELECT cafe, drink, price * 120 AS priceInYen  
FROM Sells;
```



Complex Conditions in WHERE Clause

- From Sells (cafe, drink, price), find the price **Caffe bene** charges for Mocha:

```
SELECT price
FROM Sells
WHERE cafe = 'Caffe bene' AND
      drink = 'Mocha';
```



WHERE Clause Syntax

What you can use in WHERE:

- attribute names of the relation(s) used in the FROM.
- comparison operators: =, <>, <, >, <=, >=
- apply arithmetic operations: stockprice*2
- operations, comparisons on strings
- pattern matching: s LIKE p
- special stuff for comparing dates and times.

See the textbook for more...

Then, create bigger expressions using Boolean connectives



Syntax: Patterns and “LIKE”

- WHERE clauses can have conditions in which a string is compared with a pattern, to see if it matches.
- General form: **<Attribute> LIKE <pattern>**
or **<Attribute> NOT LIKE <pattern>**
- Pattern is a quoted string with
 - **%** “any string”
 - **_** “any character.”



Example

- From Customers(name, addr, phone) find the customers that have phone numbers with middle three numbers 555:

```
SELECT name  
FROM Customers  
WHERE phone LIKE '%555- _ _ _ _';
```



Ordering the results

```
SELECT *  
FROM Account  
WHERE Name LIKE 'J%'  
ORDER BY Balance
```

- Ordering is ascending, unless you specify the **DESC** keyword.
- Ties are broken by the second attribute on the ORDER BY list, etc.



Summary

- A data model is an **abstract model** that **organizes elements of data** and standardizes how they **relate to one another** and to properties of the real world entities [wikipedia].
- There are several types of data models: Network, Hierarchal,...
 - We will study the relational, NoSQL and graph models.
- Relational database model
 - Data is structured as relations
 - Defines a limited set of operations (query and modification) to interact with the data
- Structured Query Language (SQL) is a declarative language (standard) that enable users/applications to query and modify the database.