



# MongoDB and Map-Reduce

**Abdu Alawini**

University of Illinois at Urbana-Champaign

CS411: Database Systems

**June 28, 2020**

Some slides were adopted from D. Maier. S. Davidson with permission

© 2020 A. Alawini



# Learning Objectives

After this lecture, you should be able to:

- Write MongoDB cursor queries.
- Query documents by reference.
- Write aggregation queries
- Write map-reduce programs in MongoDB



# Cursor methods

- .count, .pretty, .sort etc are all examples of cursor methods
- **.toArray** returns an array that contains all documents from a cursor

```
nsf= db.awards.find({"by": "National Science Foundation").toArray()  
if (nsf.length >0) {printjson (nsf[0])}
```

- **.forEach** applies a JavaScript function to each document from the cursor (similar to .map)

```
db.awards.find(). forEach  
( function(myDoc) { printjson ("Award: " + myDoc.name); });
```



# Joins in MongoDB

## 1. Embedded Relationships

- “Do joins while write, not on reads.”

## 2. Referenced Relationships

- Use semi-joins to get an array of keys from the first collection on which to search the second collection for matches using cursor methods



# Relationships: Embedded

```
{ _id: 1,  
  name: { first: "John", last: "Backus" },  
  birthyear: 1924,  
  contribs: [ "Fortran", "ALGOL",  
              "Backus Naur Form", "FP" ],  
  awards: [ {title: "National Medal of Science",  
             by: "National Science Foundation",  
             year: 1975 },  
            {title: "Turing Award",  
             by: "ACM",  
             year: 1977},  
            {title: "Career Award",  
             by: "NIH",  
             year: 1980},  
            {title: "Career Success Award",  
             by: "NASA",  
             year: 1982} ] }
```



## People:

## Relationships: Referenced

```
{ _id: 1,  
  name: { first: "John", last: "Backus" },  
  birthyear: 1924,  
  contribs: [ "Fortran", "ALGOL",  
              "Backus Naur Form", "FP" ],  
  awards: [ { award_id: "NMS001", year: 1975 },  
            { award_id: "TA99", year: 1977} ] }
```

## Awards:

```
{_id: "NMS001",  
 title: "National Medal of Science",  
 by: "National Science Foundation"},  
{_id: "TA99",  
 title: "Turing Award",  
 by: "ACM" }
```



## “Semijoins”

- Suppose you want to print people who have won Turing Awards using referenced relationship
  - Problem: object id of Turing Award is in collection “awards”, collection “people” references it.

```
turing= db.awards.findOne({title: "Turing Award"})  
db.people.find({"awards.award_id": turing._id})
```

- But this only works for one award with title “Turing Award”, suppose there were more.



## Awards:

```
{_id: "NMS001",  
  title: "National Medal of Science",  
  by: "National Science Foundation"},  
{_id: "TA99",  
  title: "Turing Award",  
  by: "ACM"},  
{_id: "CA85",  
  title: "Career Award",  
  by: "NIH",  
  year: 1980},  
{_id: "CSA19",  
  title: "Career Success Award",  
  by: "NASA",  
  year: 1982}
```





## Iterating using cursors

- Now suppose we want to find all people who won awards with title that starts with 'ca'

```
let ca_award= db.awards.find ({title: /^ca/i},{_id:1})
let award_Ids = new Array();
while (ca_award.hasNext()) {
    let tmp = award_Ids .push(ca_award.next()._id)
}
db.people.find({"awards.award_id":{$in: award_Ids })
```



# Aggregation

- A framework to provide “group-by” and aggregate functionality without the overhead of map-reduce.
- Conceptually, documents from a collection pass through an aggregation pipeline, which transforms the objects as they pass through (similar to UNIX pipe “|”)
- Operators include: \$project, \$match, \$group, \$sort, \$skip, \$limit, \$unwind

# Aggregation Example\*

Collection  
↓  
db.orders.aggregate( [  
 \$match stage → { \$match: { status: "A" } },  
 \$group stage → { \$group: { \_id: "\$cust\_id", total: { \$sum: "\$amount" } } }  
] )

{ cust_id: "A123", amount: 500, status: "A" }
{ cust_id: "A123", amount: 250, status: "A" }
{ cust_id: "B212", amount: 200, status: "A" }
{ cust_id: "A123", amount: 300, status: "D" }

orders

\$match →

{ cust_id: "A123", amount: 500, status: "A" }
{ cust_id: "A123", amount: 250, status: "A" }
{ cust_id: "B212", amount: 200, status: "A" }

\$group →

Results	
{	<code>_id: "A123",</code>
	<code>total: 750</code>
}	
<hr/>	
{	<code>_id: "B212",</code>
	<code>total: 200</code>
}	

\*<https://docs.mongodb.com/manual/aggregation/>



## Aggregation: \$group

- Every group expression must specify an `_id` field.
- Suppose we wanted to find how many people were born each year

```
> db.people.aggregate( { $group :  
    { _id : "$birthyear", birthsPerYear : { $sum : 1 } } })
```

```
{ "_id" : 1924, "birthsPerYear" : 1 }
```

- Contrast with aggregate operation over entire result

```
> db.people.count( )  
> db.people.find( { "name.first" : "John" }).count( )  
> db.people.count( { "name.first" : "John" })
```



## Aggregation: \$unwind

- Deconstructs an array field to output a document for each element.

```
Posts: {  
  _id : ObjectId("4c4ba5co672c685e5e8aabf3"),  
  author : "Kevin",  
  date : new Date("February 2, 2012"),  
  text : "About MongoDB...",  
  birthyear: 1980,  
  tags : [ "tech", "databases" ]  
}
```

```
>db.posts.aggregate( { $project : { author : 1, tags : 1 } }, { $unwind : "$tags" } )
```



## Result of unwind

```
>db.posts.aggregate( { $project : { author : 1, tags : 1 } }, { $unwind : "$tags" } )
```

```
{ "_id" : ObjectId("4c4ba5c0672c685e5e8aabbf3"),  
  "author" : "Kevin",  
  "tags" : "tech" },  
{ "_id" : ObjectId("4c4ba5c0672c685e5e8aabbf3"),  
  "author" : "Kevin",  
  "tags" : "databases" }
```



## Map-Reduce: Motivation

- SQL is a great way of doing batch (bulk) operations to collections (tables)
- But it's somewhat limited in the kinds of operations it supports (by default) –  
SELECT ... FROM ... WHERE ... GROUP BY ... HAVING ...
- What if we want to do *arbitrary* selection predicates (in Java etc.), and arbitrary aggregations (in Java etc.)
  - User defined functions, or alternative frameworks
- What if we wanted to scale this up to run on a 10000 node compute cluster??



# Map-Reduce Analogy: National census

- Suppose we have 10,000 employees, whose job is to collect census forms and to determine how many people live in each city
- How would you organize this task?

The image shows a sample of the 2010 U.S. Census form, Form D-61, titled "U.S. Census 2010". The form is from the U.S. Department of Commerce, Economics and Statistics Administration, U.S. Census Bureau. It includes instructions for users to use a blue or black pen and to start at the top. The form contains several sections with questions and checkboxes for demographic data collection, including questions about the number of people living in the household, the sex of the person, the date of birth, the race, and the telephone number. The form is labeled "Form D-61 (1-15-2009)" and "USCENSUSBUREAU".

[http://www.census.gov/2010census/pdf/2010\\_Questionnaire\\_Info.pdf](http://www.census.gov/2010census/pdf/2010_Questionnaire_Info.pdf)

16





## Making things more complicated

- Suppose workers take vacations, get sick, work at different rates
- Suppose some forms are incorrectly filled out and require corrections or need to be thrown away
- What if the supervisor gets sick?
- How big should the stacks be?
- How do we monitor progress?
- ...




# I don't want to deal with all this!!!

- Wouldn't it be nice if there were some system that took care of all these details for you?
- Ideally, you'd just tell the system what needs to be done
- That's the MapReduce framework.



# Abstracting into a digital data flow





## Abstracting once more

- There are two kinds of workers:
  - Those who take input data items and produce output items for the “stacks”
  - Those who take the stacks and **aggregate** the results to produce outputs on a per-stack basis

map does SELECT + PROJECT

- We'll call these:

- **map**: takes (item\_key, value), produces one or more (stack\_key, value') pairs
- **reduce**: takes (stack\_key, {set of value'}), produces one or more output results – typically (stack\_key, agg\_value)

We will refer to this key  
as the reduce key; it is like the GROUP BY key



## Why MapReduce?

- Scenario:
  - You have a huge amount of data, e.g., all the Google searches of the last three years
  - You would like to perform a computation on the data, e.g., find out which search terms were the most popular
- Analogy to the census example:
  - The computation isn't necessarily difficult, but parallelizing and distributing it, as well as handling faults, is challenging
- Idea: A programming abstraction / template!
  - Write a simple program to express the (simple) computation, and let the language runtime do all the hard work



## Simple example: Word count

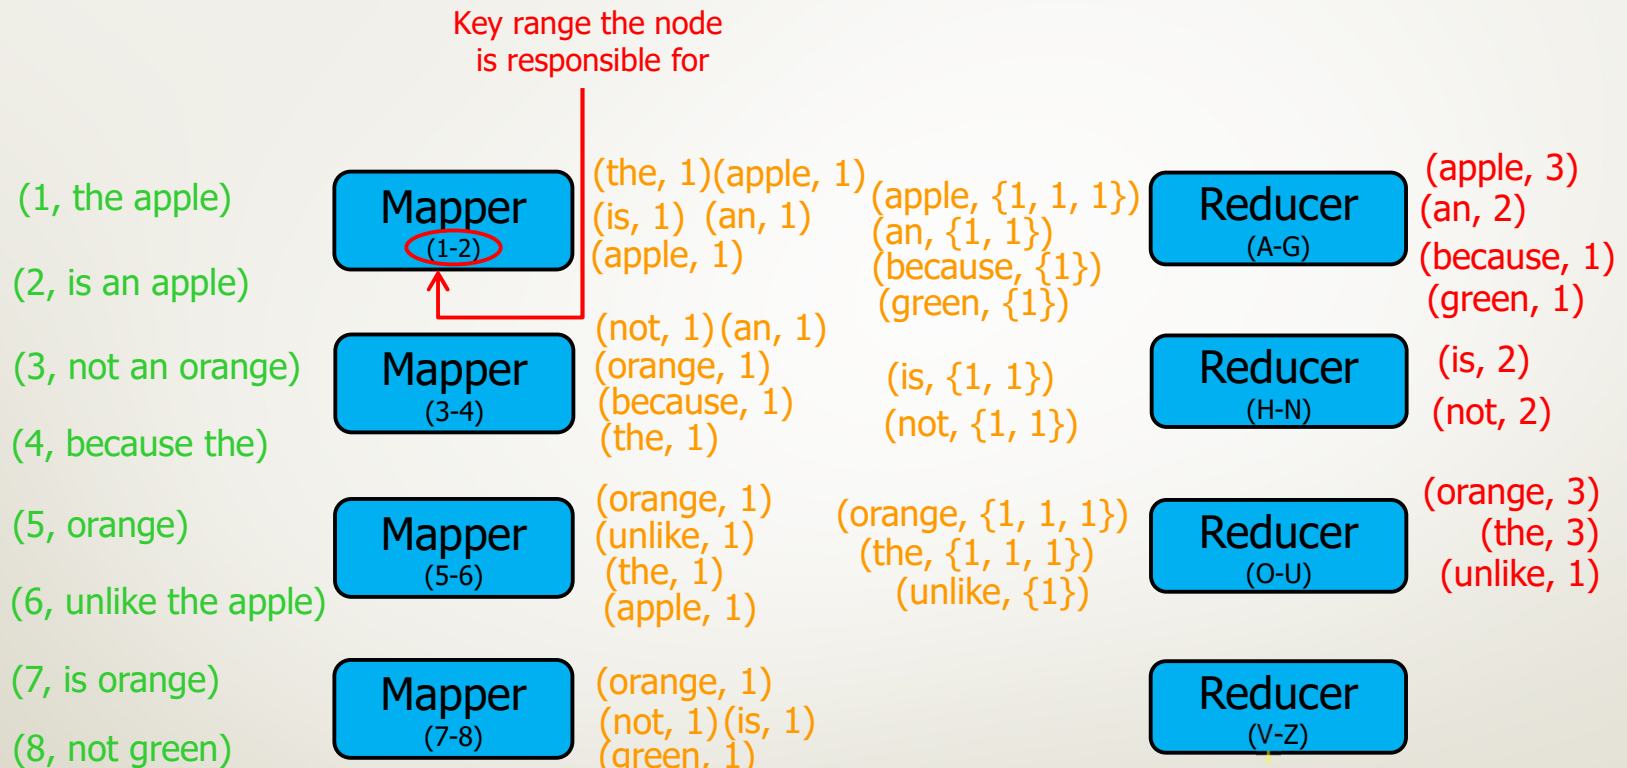
```
map(String key, String value) {  
    // key: document name, line no  
    // value: contents of line  
    for each word w in value:  
        emit(w, "1")  
}
```

```
reduce(String key, Iterator values)  
{  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    emit(key, result)  
}
```

- **Goal:** Given a set of documents, count how often each word occurs
  - Input: Key-value pairs (document:lineNumber, text)
  - Output: Key-value pairs (word, #occurrences)
  - What should be the intermediate key-value pairs?



# Simple example: Word count



①

Each mapper receives some of the KV-pairs as input

②

The mappers process the KV-pairs one by one

③

Each KV-pair output by the mapper is sent to the reducer that is responsible for it

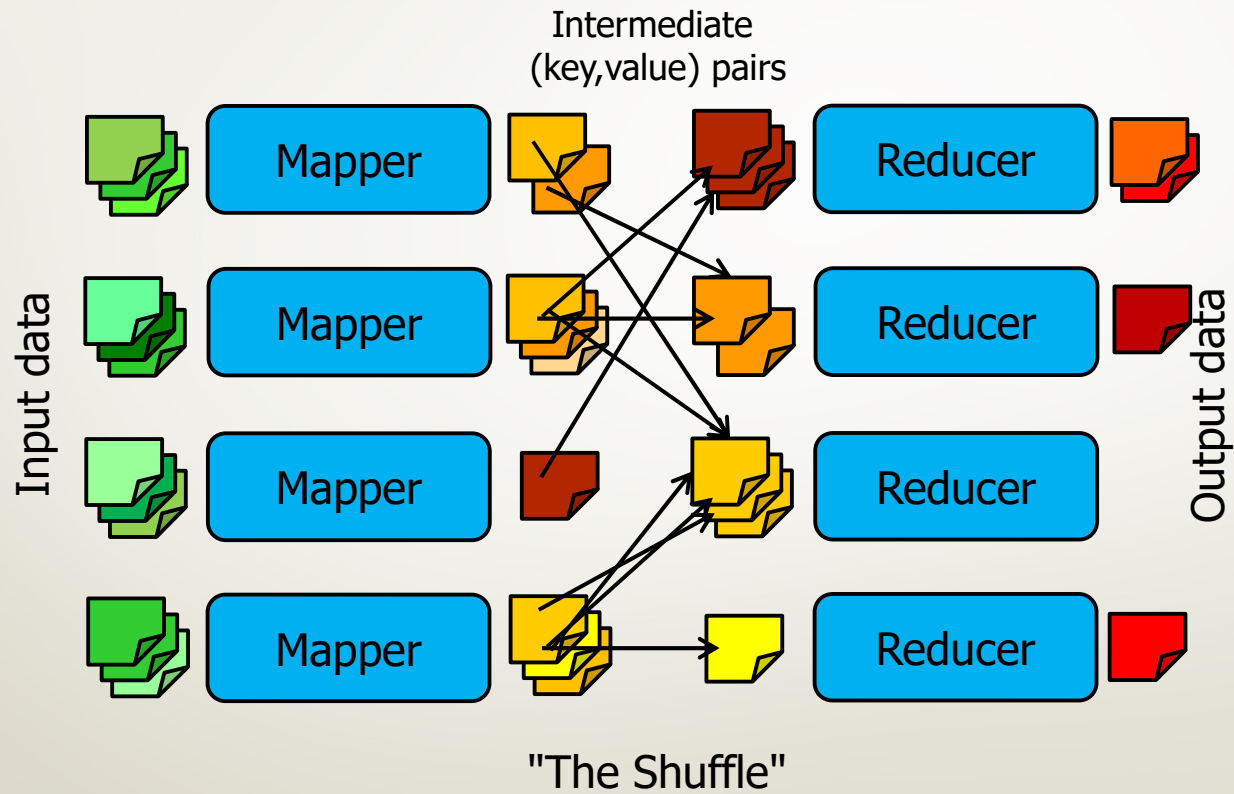
④

The reducers sort their input by key and group it

⑤

The reducers process their input one group at a time

## MapReduce dataflow







## More examples

- Count URL access frequency
  - Map: output each URL as key, with count 1
  - Reduce: sum the counts
- Reverse web-link graph
  - Map: output (target, source) pairs when link to target found in source
  - Reduce: concatenates values and emits (target, list(source))
- Inverted index
  - Map: Emits (word, documentID)
  - Reduce: Combines these into (word, list(documentID))



# Designing MapReduce algorithms

- Key decision: What should be done by `map`, and what by `reduce`?
  - `map` can do something to each individual key-value pair, but it can't look at other key-value pairs
    - Example: Filtering out key-value pairs we don't need
  - `map` can emit more than one intermediate key-value pair for each incoming key-value pair
    - Example: Incoming data is text, `map` produces `(word,1)` for each word
  - `reduce` can aggregate data; it can look at multiple values, as long as `map` has mapped them to the same (intermediate) key
    - Example: Count the number of words, add up the total cost, ...
- Need to get the intermediate format right!
  - If `reduce` needs to look at several values together, `map` must emit them using the same key!



# Common mistakes to avoid

- Mapper and reducer should be **stateless**
  - Don't use static variables - after map + reduce return, they should remember nothing about the processed data!
  - Reason: No guarantees about which key-value pairs will be processed by which workers!
- Don't try to do your own **I/O!**
  - Don't try to write to files in the file system
  - The MapReduce framework does the I/O for you (usually):
    - All the incoming data will be fed as arguments to map and reduce
    - Any data your functions produce should be output via emit

```
HashMap h = new HashMap();  
map(key, value) {  
    if (h.contains(key)) {  
        h.add(key, value);  
        emit(key, "X");  
    }  
}
```

**Wrong!**

```
map(key, value) {  
    File foo =  
        new File("xyz.txt");  
    while (true) {  
        s = foo.readLine();  
        ...  
    }  
}
```

**Wrong!**



## More common mistakes to avoid

```
map(key, value) {  
    emit("FOO", key + " " + value);  
}
```

Wrong!

```
reduce(key, value[]) {  
    /* do some computation on  
    all the values */  
}
```

- Mapper must not map too much data to the same key
  - In particular, don't map *everything* to the same key!! Otherwise the reduce worker will be overwhelmed!
  - It's okay if some reduce workers have more work than others
    - Example: In WordCount, the reduce worker that works on the key 'and' has a lot more work than the reduce worker that works on 'syzygy'.



# MapReduce and MongoDB

- Suppose we have a collection of orders that look like the following:

```
{  
  _id: "o123",  
  cust_id: "abc123",  
  ord_dat: new Date("Oct 04, 2013"),  
  status: 'A',  
  price: 17,  
  items: [ {sku: "mmm", qty: 5, price: 2.5},  
            {sku: "nnn", qty: 3, price: 1.5} ]  
}
```



## Example 1: MapReduce and MongoDB

- Return the total cost per customer.

```
var mapFunction1 = function() {  
    emit(this.cust_id, this.price);  
}  
  
var reduceFunction1 = function(keyCustId, valuesPrices) {  
    return Array.sum(valuesPrices);  
}  
  
db.orders.mapReduce(mapFunction1, reduceFunction1,  
    {out: {inline:1}})
```



## Output on the following input...

input

```
order1= {_id: "o123", cust_id: "abc123",..., price: 17, ...}  
order2= {_id: "o124", cust_id: "abc124",..., price: 20, ...}  
order3= {_id: "o125", cust_id: "abc123", ..., price: 13,...}
```

output

```
{"results": [{"_id": "abc123", "value": 30},  
             {"_id": "abc124", "value": 20}],  
"timeMillis": 23,  
"counts": {"input": 3, "emit": 3, "reduce": 1, "output": 2},  
"ok": 1}
```



## Example 2: MapReduce and MongoDB

- Return the total and average number of each item (SKU) in orders with status “A”.





## Example 2: MapReduce and MongoDB

- Return the total and average number of each item (SKU) in orders with status “A”.

```
var mapFunction2 = function () {  
    for (var idx=0; idx<this.items.length; idx++) {  
        var key= this.items[idx].sku;  
        var value= { count: 1, qty: this.items[idx].qty };  
        emit(key,value); } }  
  
var reduceFunction2= function(keySKU, countObjVals) {  
    reducedVal= {count: 0, qty: 0};  
    for (var idx=0; idx< countObjVals.length; idx++) {  
        reducedVal.count += countObjVals[idx].count;  
        reducedVal.qty += countObjVals[idx].qty; }  
    return reducedVal; }
```



## Example 2, cont

- But we need more to calculate the average, hence have a “finalize” function.

```
var finalizeFunction2= function (key, reducedVal) {  
    reducedVal.avg= reducedVal.qty/reducedVal.count;  
    return reducedVal;  
};
```

- Putting it all together:

```
db.orders.mapReduce(mapFunction2, reduceFunction2,  
    {out:{inline:1},  
    query: {status: "A"},  
    finalize: finalizeFunction2 } )
```



## Output (Ex.2)

```
{"results" : [{"_id" : "mmm",  
  "value" : {"count" : 3, "qty" : 14,  
    "avg" : 4.666666666666667}},  
  {"_id" : "nnn",  
    "value" : {"count" : 3, "qty" : 13,  
      "avg" : 4.333333333333333} } ],  
  "timeMillis" : 13,  
  "counts" : {"input" : 3,  
    "emit" : 6,  
    "reduce" : 2,  
    "output" : 2},  
  "ok" : 1 }
```



# Conclusions

- Big data is two problems: analysis and storage
  - **Analysis:** relies on modeling, ML, statistics
  - **Storage:** relies on database technique to store and manipulate huge amounts of data to facilitate fast queries
- Due to “variety” and “volume”, relational solutions frequently fall short
  - Need for less structure and nested information
  - Need for parallelism
- Map-reduce is part of many of the key-value store NoSQL solutions.