



# Map-Reduce Neo4j: Graph Database

**Abdu Alawini**

University of Illinois at Urbana-Champaign

CS411: Database Systems

**June 28, 2020**



# Learning Objectives

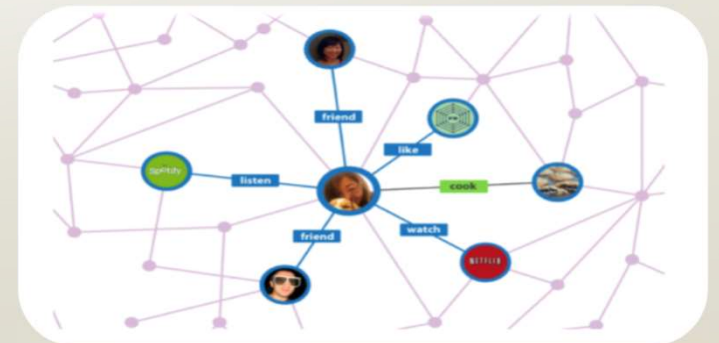
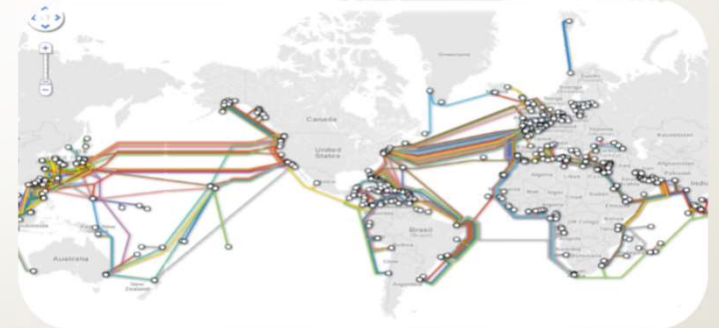
After this lecture, you should be able to:

- Discuss the Labeled Property Graph Data Model
- Write basic nodes and relationships Cypher queries.
- Use Cypher commands to manipulate Neo4J graph data.

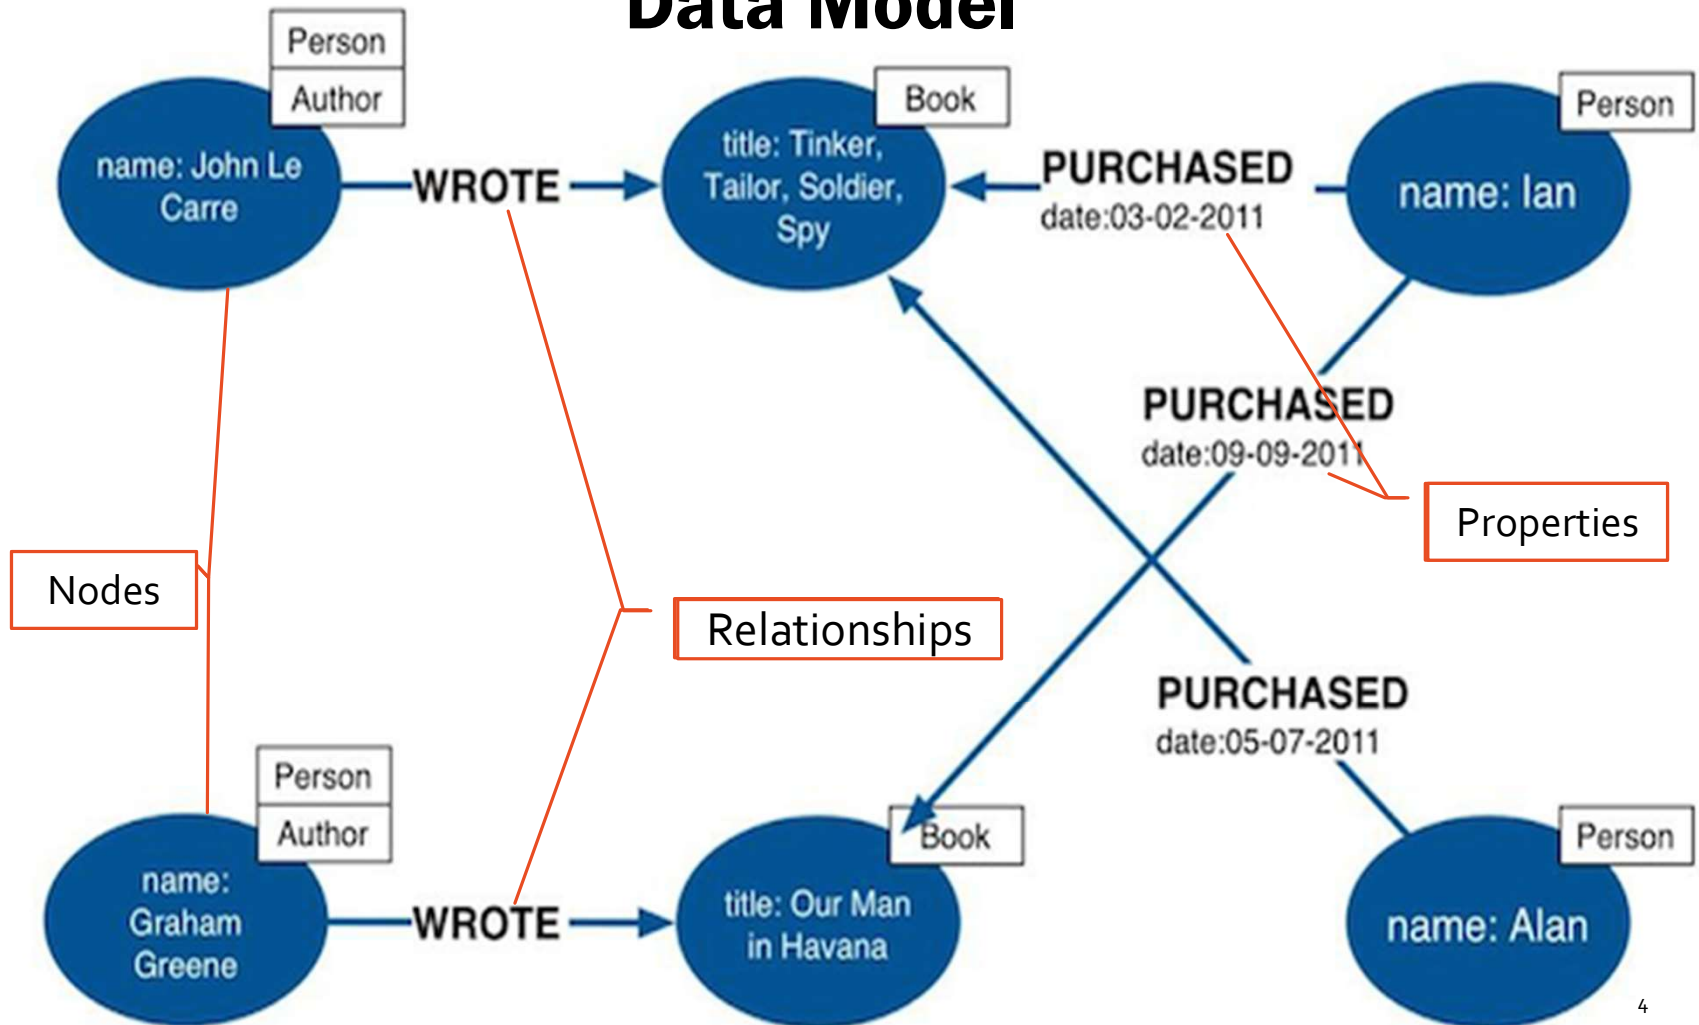


# Neo4j: Graph Database

- Many types of data can be represented as graphs
  - Road networks, with intersections as nodes and road segments as edges
  - Computer networks, with computers as nodes and connections as edges
  - Social networks, with people/postings as nodes and edges as relationship (e.g. friends, likes, created, ...)
- Graph databases store relationships and connections as first-class entities: “Property Graph Model”



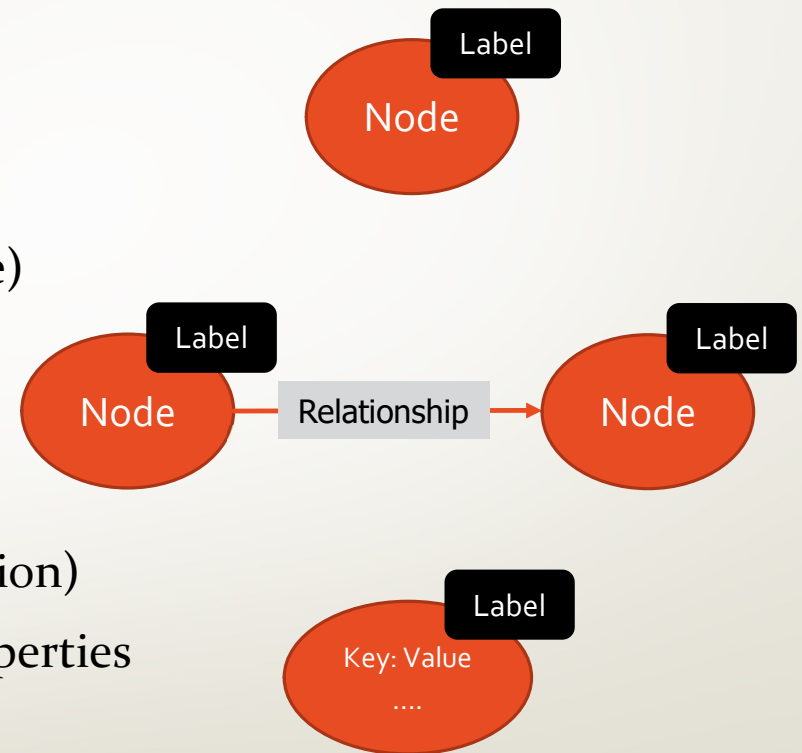
# Labeled Property Graph Data Model





# Neo4j Nodes and Relationships

- Nodes
  - have a system-assigned id
  - can have key/value properties
  - there is a reference node (“starting point” into the node space)
- Relationships (Edges)
  - have a system-assigned id
  - are directed (but can be traversed in either direction)
  - have a type – can have key/value properties
- Key/value properties
  - values always stored as strings
  - support for basic types and arrays of basic types





# The Power of Graph Databases

- **Performance**

- Graph databases have better performance when dealing with connected data vs. relational databases and NOSQL

- **Flexibility**

- No need to define schema upfront

- **Agility**

- graph data model is schema-free
- testable graph database's API
- query language

} Graph DB and application  
evolve in an agile fashion



# Cypher: A Graph Query Language

- SQL-like syntax
- Declarative pattern-matching graph query language
- Query a graph DB to find data (Nodes, Relationships, subgraphs) that matches a specific pattern
  - uses ASCII to specify a patterns

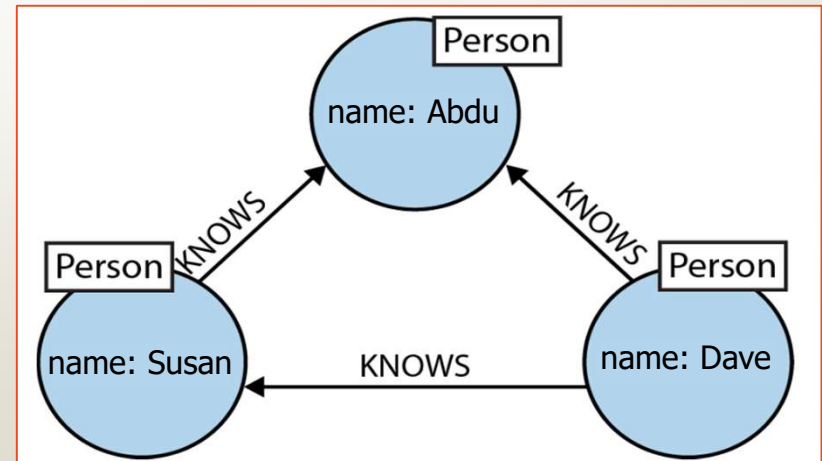
# Cypher Basics

## Variables (Identifiers)

- **Patterns** describes *path* that connects a node (or set of nodes)
- **Identifiers** allow us to refer to the same node more than once when describing a pattern
- In Cypher, graphs are described using *specification by example*.

## Cypher Pattern

```
(x:Person {name: 'Abdu'})  
<-[:KNOWS]-(y:Person {name: 'Dave'})  
-[:KNOWS]->(z:Person {name: 'Susan'})  
-[:KNOWS]->(x)
```







# Cypher Queries

- Cypher is comprised of three main concepts
  - **MATCH**: graph pattern to match, bound to the starting points
  - **WHERE**: filtering criteria
  - **RETURN**: what to return
- Implemented using the Scala programming language



# Querying Nodes

- Nodes are surrounded with parentheses

```
()  
(matrix)  
(:Movie)  
(matrix:Movie)  
(matrix:Movie {title: "The Matrix"})  
(matrix:Movie {title: "The Matrix", released:  
1997})
```

- Nodes queries

```
MATCH (node:Label)  
WHERE node.property = {value}  
RETURN node.property
```



# Examples

- Find "Apollo 13" movie

```
MATCH (n:Movie {title:"Apollo 13"})  
RETURN n.released;
```

n.released

1995

- Find 1990's movies

```
MATCH (n:Movie)  
WHERE n.released >= 1990 AND n.released <=2000  
RETURN n;
```

| "n"  |
|--|
| {"tagline":"Welcome to the Real World","title":"The Matrix", "released":"1999"}  |
| {"tagline":"At odds in life... in love on-line.", "title":"When Harry Met Sally", "released":"1998"}   |
| {"tagline":"In every life there comes a time when that thing you dream becomes that thing you do", "title":"That Thing You Do", "released":"1996"} |
| {"tagline":"Pain heals, Chicks dig scars... Glory lasts forever", "title":"The Replacements", "released":"2000"}                                   |



# Querying Relationships

Relationships are basically an arrow --> between two nodes.

- relationship-types `-[:KNOWS]->` , `<-[: LIKE]-`
- a variable name `-[rel:KNOWS]->` before the colon
- additional properties `-[rel:KNOWS {since:2018}]->`
- structural information for paths of variable length `-[:KNOWS*..4]->`

```
MATCH (n1:Label1)-[rel:TYPE]->(n2:Label2)
WHERE rel.property = {value}
RETURN rel.property, type(rel)
```

## Example

```
MATCH (m:Movie)<-[r:ACTED_IN]-(p:Person)
WHERE r.roles = "Morpheus"
RETURN p.name, m.title;
```

| "p.name"             | "m.title"                |
|----------------------|--------------------------|
| "Laurence Fishburne" | "The Matrix"             |
| "Laurence Fishburne" | "The Matrix Reloaded"    |
| "Laurence Fishburne" | "The Matrix Revolutions" |



# Patterns

- Nodes and relationship expressions are the building blocks for more complex patterns.
- Patterns can be written continuously or separated with commas.
- You can refer to variables declared earlier or introduce new ones.

## Types of Pattern Matching

- friend-of-a-friend

```
(user)-[:KNOWS]-(friend)-[:KNOWS]-(foaf)
```

- shortest path:

```
path = shortestPath( (user)-[:KNOWS*..5]-(other) )
```

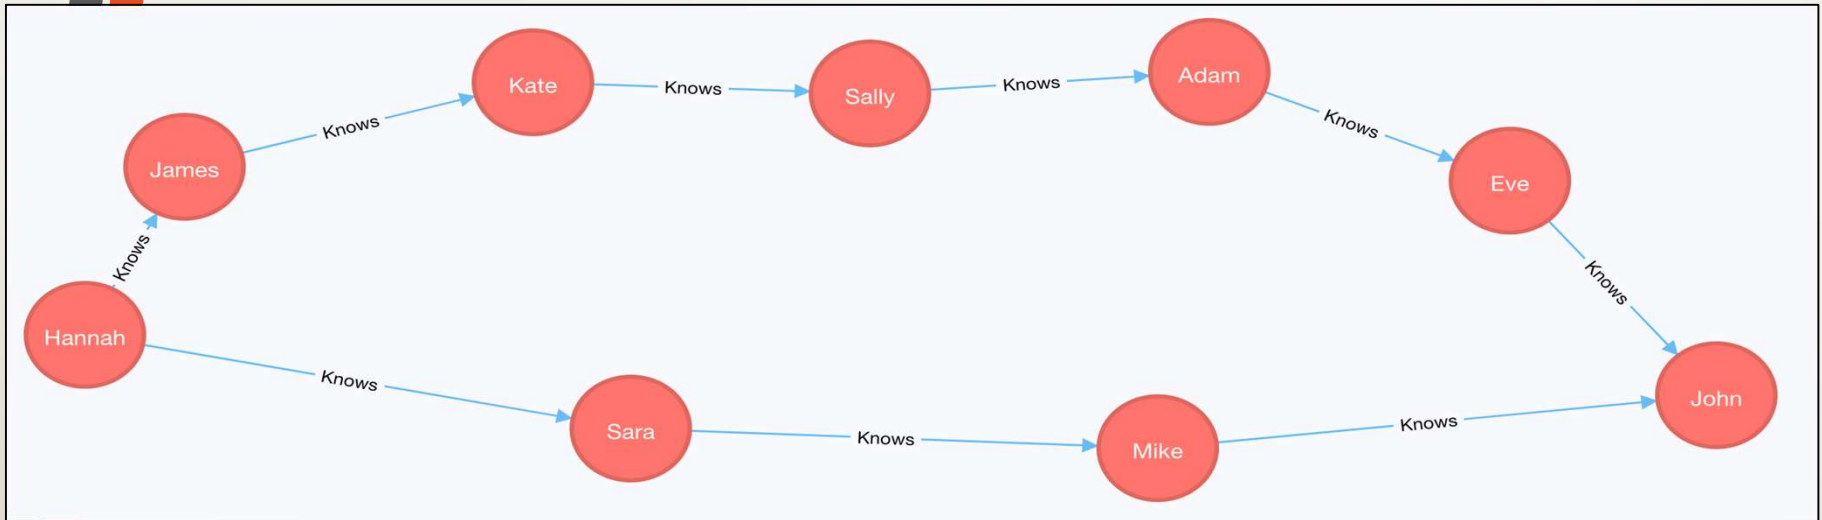
- collaborative filtering

```
(user)-[:PURCHASED]->(product)<-[:PURCHASED]-()-[:PURCHASED] ->(otherProduct)
```

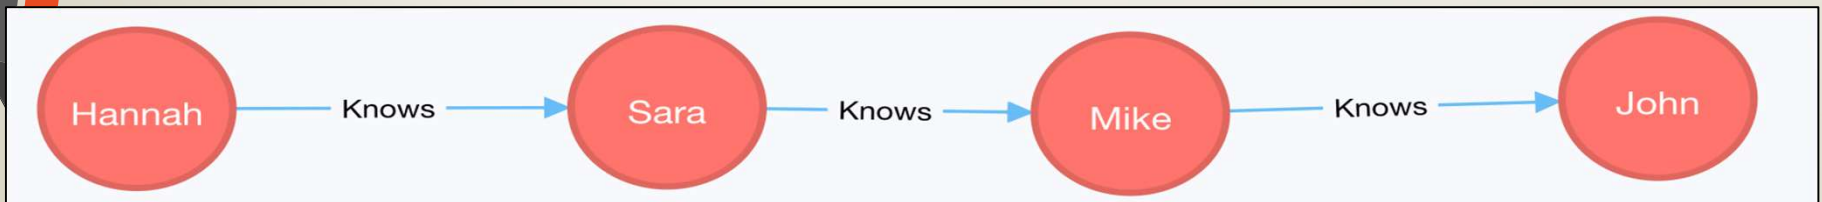
- tree navigation

```
(root)<-[:PARENT*]-(leaf:Category)-[:ITEM]->(data:Product)
```

# Shortest path example



```
MATCH x = shortestPath( (s:Student {name:"Hannah"})-[:Knows*..10]->(s1:Student {name:"John"}))  
RETURN x;
```





# Collaborative Filtering Example

```
MATCH (p:Person)-[:ACTED_IN]->(m:Movie)<-[:DIRECTED]-(p)
RETURN p.name, m.title;
```

| p.name         | m.title           |
|----------------|-------------------|
| Tom Hanks      | That Thing You Do |
| Clint Eastwood | Unforgiven        |
| Danny DeVito   | Hoffa             |





# Outline

- ✓ Map-reduce
- ✓ Introduction to Neo4J
- Cypher: A Graph Query Language
  - ✓ Querying Nodes and Relationships using Patterns
  - Manipulating Graph Data



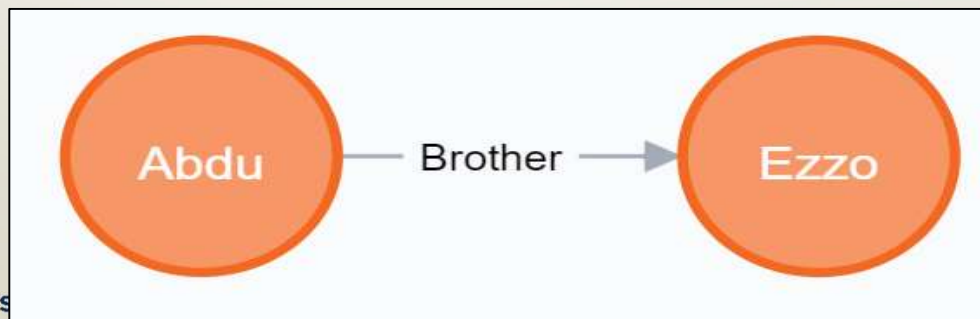
# Manipulating Graph Data

**Create** statement creates nodes and relationships specified in the pattern

```
1 CREATE (var:Person {name:"Abdu"})  
2 RETURN var;
```



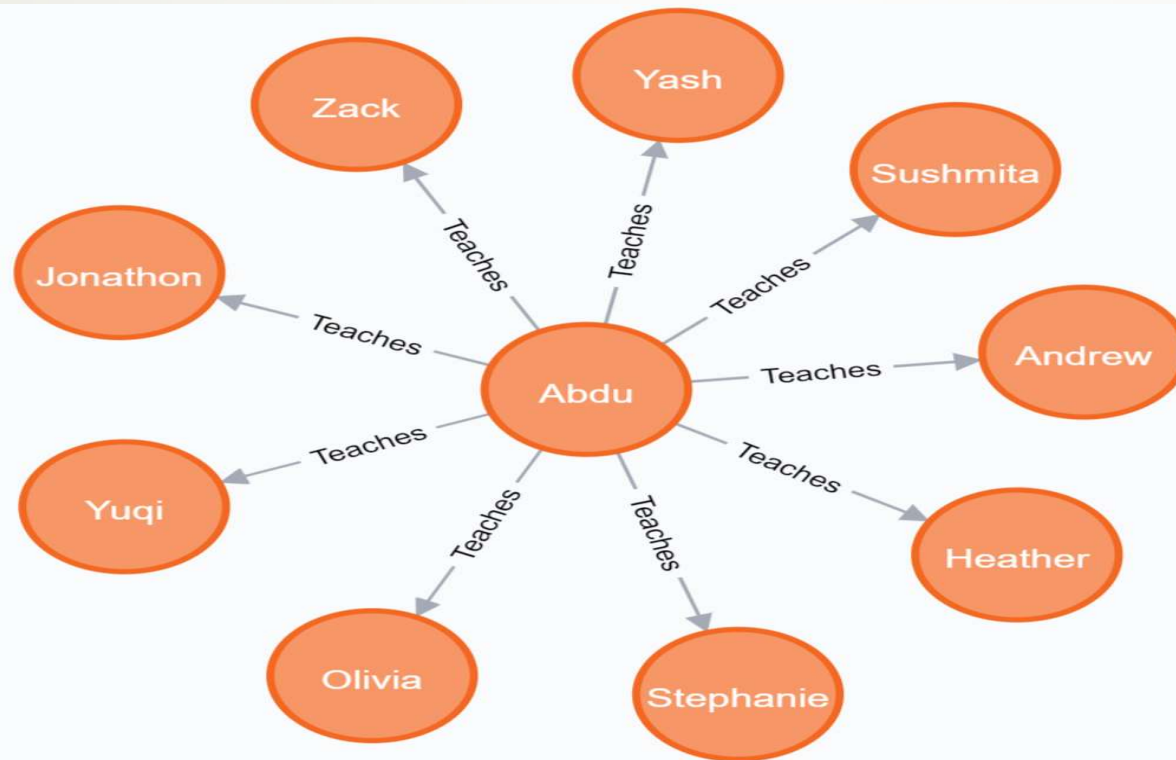
```
1 CREATE path = (:Person {name:"Abdu", age:21})-[:Brother]->(:Person {name:"Ezzo", age:55})  
2 RETURN path;
```





## FOREACH Statement

```
1 CREATE(me:Person {name: 'Abdu'})
2 FOREACH (student in ['Jonathon','Zack','Yash', 'Stephanie', 'Olivia',
  'Sushmita','Heather','Yuqi','Andrew'] | CREATE (me)-[:Teaches]->(:Person {name:student}))
```



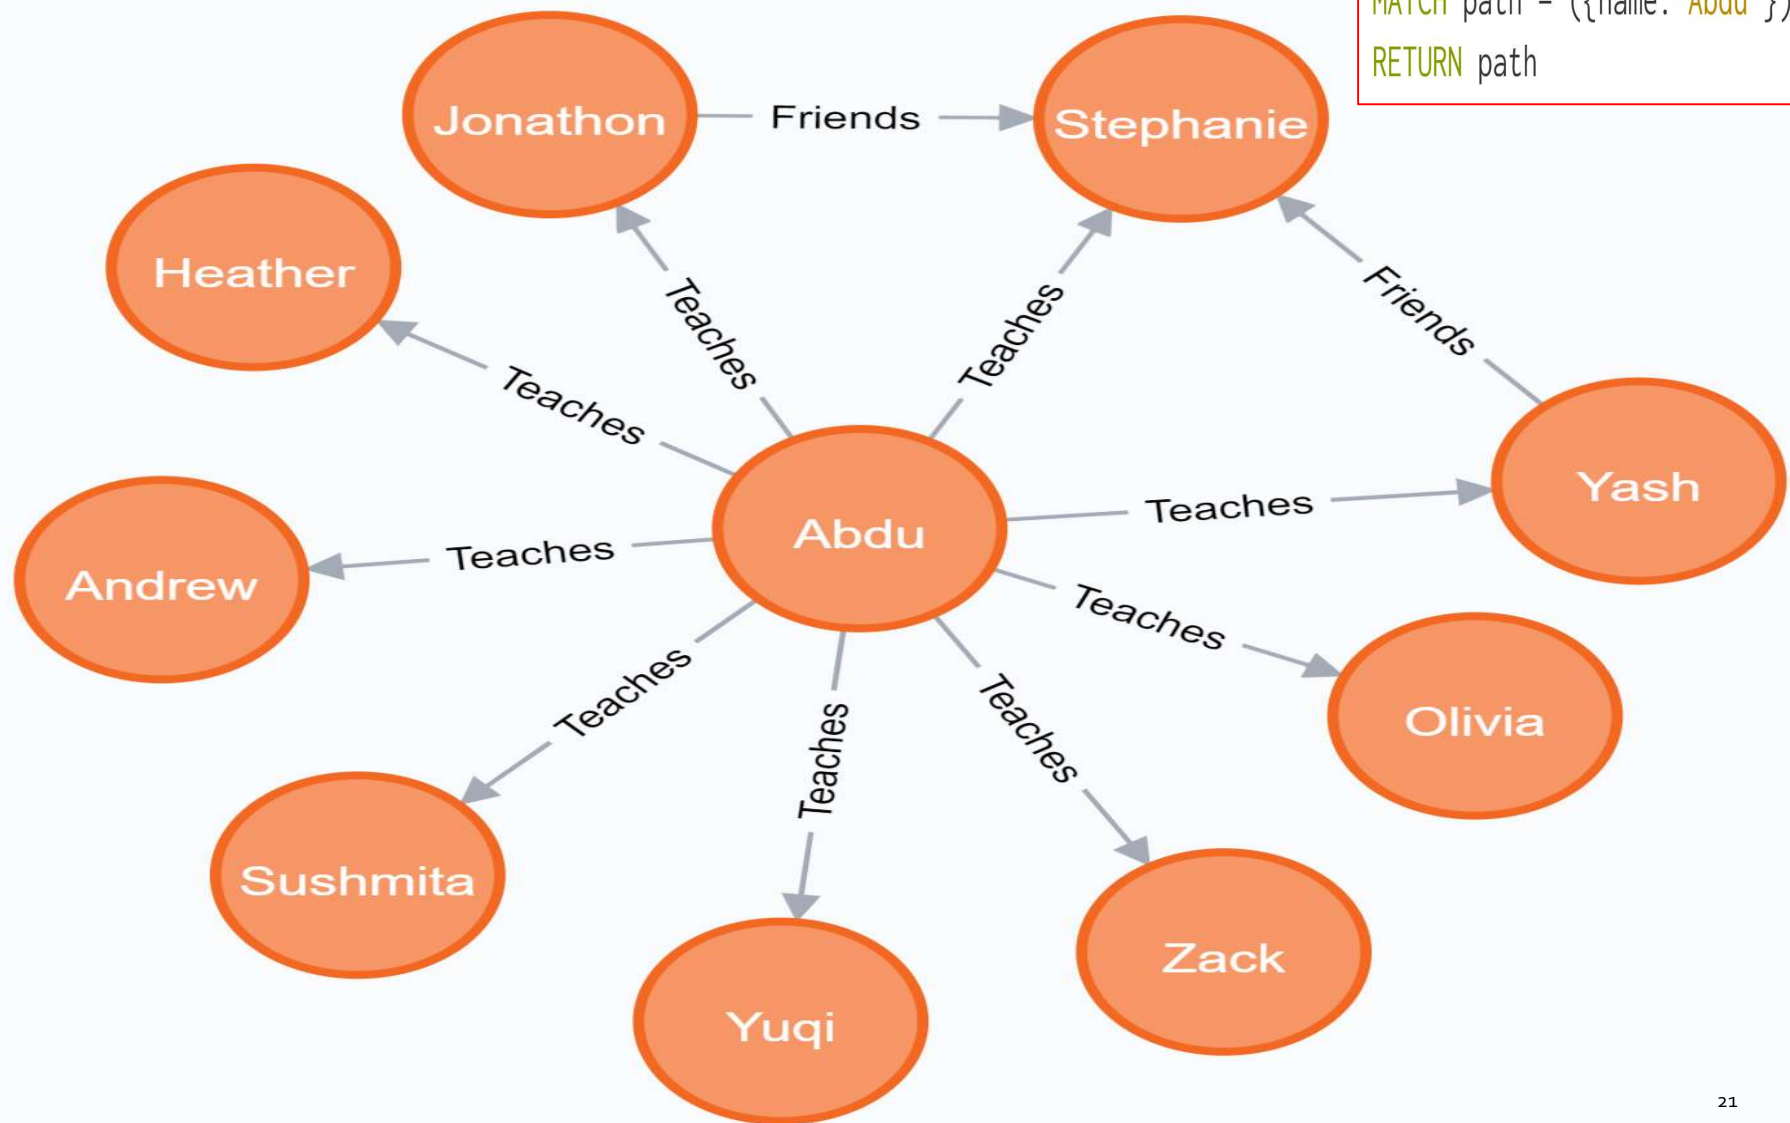
# I

## Creating Relationships

```
MATCH (a:Person {name:"Lamy"})  
MATCH (b:Person {name:"Devin"})  
MATCH (c:Person {name:"Vidisha"})  
CREATE (a)-[:Friends]->(b)<-[:Friends]-(c)  
RETURN a,b,c
```



```
MATCH path = ({name:"Abdu"})--()  
RETURN path
```





# DELETE, REMOVE and SET Commands

To delete a node:

```
MATCH (p:Person {name:"Abdu"})  
DELETE p;
```

To delete a node and  
all its relationships:

```
MATCH (a {name:"Abdu"})  
DETACH DELETE a;
```

To remove a property:

```
MATCH (abdu { name: 'Abdu' })  
REMOVE abdu.age  
RETURN abdu
```

To update or add  
a property:

```
MATCH (n { name: 'Abdu' })  
SET n.name = 'Abdussalam'  
RETURN n
```

Update a property  
using FOREACH

```
MATCH p =(begin)-[*]->(END )  
WHERE begin.name = 'A' AND END.name = 'D'  
FOREACH (n IN nodes(p)| SET n.marked = TRUE )
```



## Completing patterns: Merge

- acts like a combination of MATCH or CREATE
- checks for the existence of data first before creating it

```
MATCH (m:Person {name:"Abdu"})  
MERGE (m)-[:Works_with]->(s:Person {name:"Susan"})  
RETURN m,s
```



# I

## ON CREATE, ON MATCH

```
MERGE (a:Person { name: 'Abdu' })  
ON CREATE SET a.created = timestamp()  
ON MATCH SET a.lastSeen = timestamp()  
RETURN a.name, a.created, a.lastSeen
```

Running the above command for the first time

| "a.name" | "a.created"     | "a.lastSeen" |
|----------|-----------------|--------------|
| "Abdu"   | "1510526671422" | null         |

Running the same command for the second time

| "a.name" | "a.created"     | "a.lastSeen"    |
|----------|-----------------|-----------------|
| "Abdu"   | "1510526671422" | "1510526738428" |





# Aggregation

common aggregation functions are supported:  
count, sum, avg, min, and max

```
MATCH (p:Person)
RETURN count(*) as headcount;
```

|             |
|-------------|
| "headcount" |
| "145"       |

```
MATCH (actor:Person)-[:ACTED_IN]->(movie:Movie)<-[:DIRECTED]-(director:Person)
RETURN actor,director,count(*) AS collaborations
```

| "actor"                                    | "director"                               | "collaborations" |
|--|--|------------------|
| {"born":"1946","name":"Susan Sarandon"}    | {"born":"1965","name":"Lana Wachowski"}  | "1"              |
| {"born":"1960","name":"Annabella Sciorra"} | {"born":"1956","name":"Vincent Ward"}    | "1"              |
| {"born":"1956","name":"Tom Hanks"}         | {"born":"1951","name":"Robert Zemeckis"} | "2"              |
| {"born":"1953","name":"David Morse"}       | {"born":"1959","name":"Frank Darabont"}  | "1"              |

# I

## COLLECT

- **Collect()** function collects all aggregated values into a list

```
MATCH (m:Movie)<-[:ACTED_IN]-(a:Person)
RETURN m.title AS movie, collect(a.name) AS cast, count(*) AS actors
```

| "movie"             | "cast"  | "actors" |
|---------------------|---|----------|
| "You've Got Mail"   | ["Dave Chappelle","Parker Posey","Steve Zahn","Meg Ryan","Tom Hanks","Greg Kinnear"]                                | "6"      |
| "Apollo 13"         | ["Tom Hanks","Kevin Bacon","Ed Harris","Bill Paxton","Gary Sinise"]   | "5"      |
| "Johnny Mnemonic"   | ["Dina Meyer","Takeshi Kitano","Ice-T","Keanu Reeves"]  | "4"      |
| "Stand By Me"       | ["Marshall Bell","Kiefer Sutherland","John Cusack","Corey Feldman","Jerry O'Connell","River Phoenix","Wil Wheaton"] | "7"      |
| "The Polar Express" | ["Tom Hanks"]   | "1"      |

# I

## Composing Statements: UNION

- UNION combines the results of two statements that have the same result structure

Equivalent  
Query

```
MATCH (p:Person)-[r:ACTED_IN]->(m:Movie)
RETURN p.name as name, type(r) as Acted_Directed, m.title as title
UNION
MATCH (p:Person)-[r:DIRECTED]->(m:Movie)
RETURN p.name as name, type(r) as Acted_Directed, m.title as title

MATCH (actor:Person)-[r:ACTED_IN|DIRECTED]->(movie:Movie)
RETURN actor.name AS name, type(r) AS acted_in, movie.title AS title
```

| "name"            | "Acted_Directed" | "title"                  |
|-------------------|------------------|--------------------------|
| "Nathan Lane"     | "ACTED_IN"       | "Joe Versus the Volcano" |
| "Tom Hanks"       | "ACTED_IN"       | "Joe Versus the Volcano" |
| "Meg Ryan"        | "ACTED_IN"       | "Joe Versus the Volcano" |
| "Lilly Wachowski" | "DIRECTED"       | "The Matrix"             |
| "Lana Wachowski"  | "DIRECTED"       | "The Matrix"             |
| "Rob Reiner"      | "DIRECTED"       | "When Harry Met Sally"   |



## Composing Statements: WITH

- WITH clause combines individual parts of a query and declare which data flows from one to the other.
- WITH is like RETURN with the difference that it doesn't finish a query but prepares the input for the next part.



## WITH Example

```
MATCH (person:Person)-[:ACTED_IN]->(m:Movie)
WITH person, count(*) AS appearances, collect(m.title) AS movies
WHERE appearances > 1
RETURN person.name, appearances, movies
```

| "person.name"            | "appearances" | "movies"   |
|--------------------------|---------------|--|
| "Cuba Gooding Jr."       | "4"           | ["A Few Good Men","Jerry Maguire","As Good as It Gets","What Dreams May Come"] |
| "Oliver Platt"           | "2"           | ["Frost/Nixon","Bicentennial Man"]   |
| "Philip Seymour Hoffman" | "2"           | ["Twister","Charlie Wilson's War"]   |
| "Sam Rockwell"           | "2"           | ["The Green Mile","Frost/Nixon"]   |
| "Greg Kinnear"           | "2"           | ["As Good as It Gets","You've Got Mail"]                                       |
| "Zach Grenier"           | "2"           | ["RescueDawn","Twister"]   |
| "Rosie O'Donnell"        | "2"           | ["A League of Their Own","Sleepless in Seattle"]                               |

Source: <https://neo4j.com/developer/cypher-query-language/>



# Indexing and Constraints

- Goal of indexing: find the starting point in the graph as fast as possible

```
CREATE INDEX ON :ACTOR(name);
```

```
MATCH (p:ACTOR {name: 'Michael'}) RETURN p
```

Interested in DB Tuning?

<http://neo4j.com/docs/developer-manual/current/cypher/query-tuning/using/>

- Unique constraints guarantee uniqueness of a certain property on nodes with a specific label.

```
CREATE CONSTRAINT ON (p:Person) ASSERT p.name IS UNIQUE
```





# Index Management

- Listing database indexes

```
CALL db.indexes
```

| "description"            | "state"  | "type"                |
|--------------------------|----------|-----------------------|
| "INDEX ON :Person(name)" | "ONLINE" | "node_label_property" |

- Dropping an Index

```
DROP INDEX ON :Person(name)
```



# From Relational to Graph Databases

- Graph databases store relationships and connections as first-class entities: “Property Graph Model”

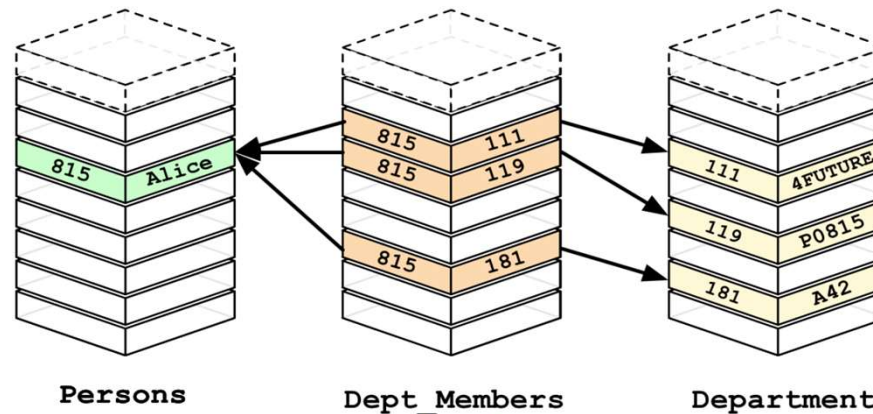
| RDBMS            | Graph Databases            |
|------------------|----------------------------|
| Tables           | Set of Nodes/Relationships |
| Rows             | Nodes                      |
| Columns and data | Data properties and values |
| Constraints      | Relationships              |
| Joins            | Traversals                 |



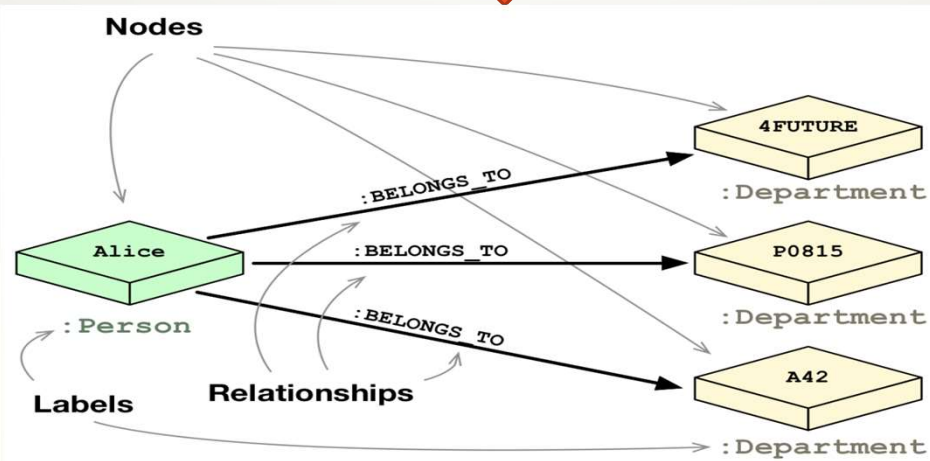


# From Relational to Graph Databases

Relational Model

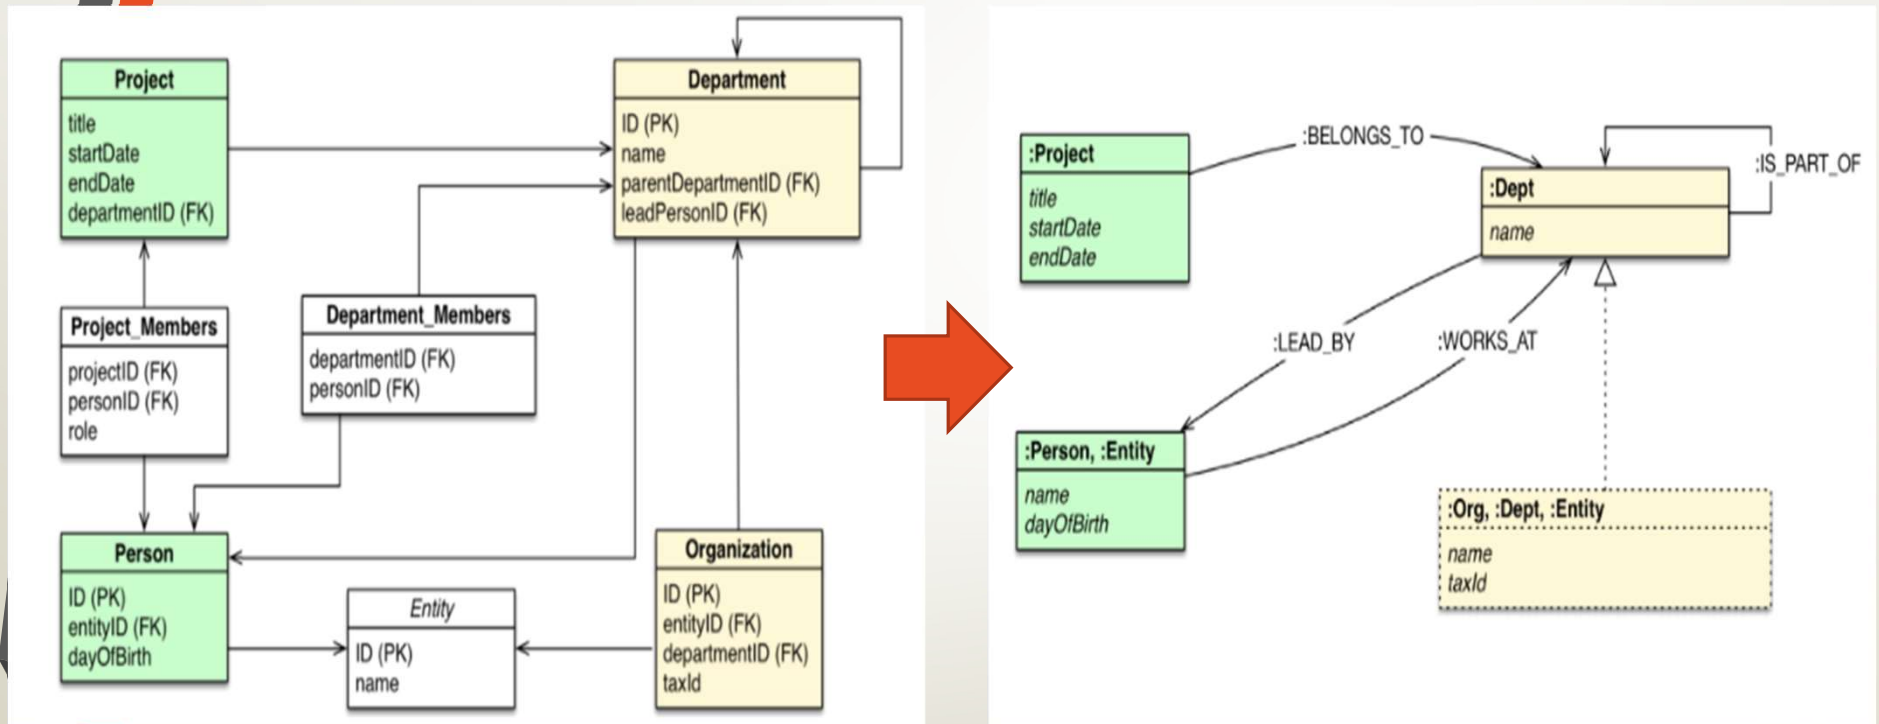


Graph Model



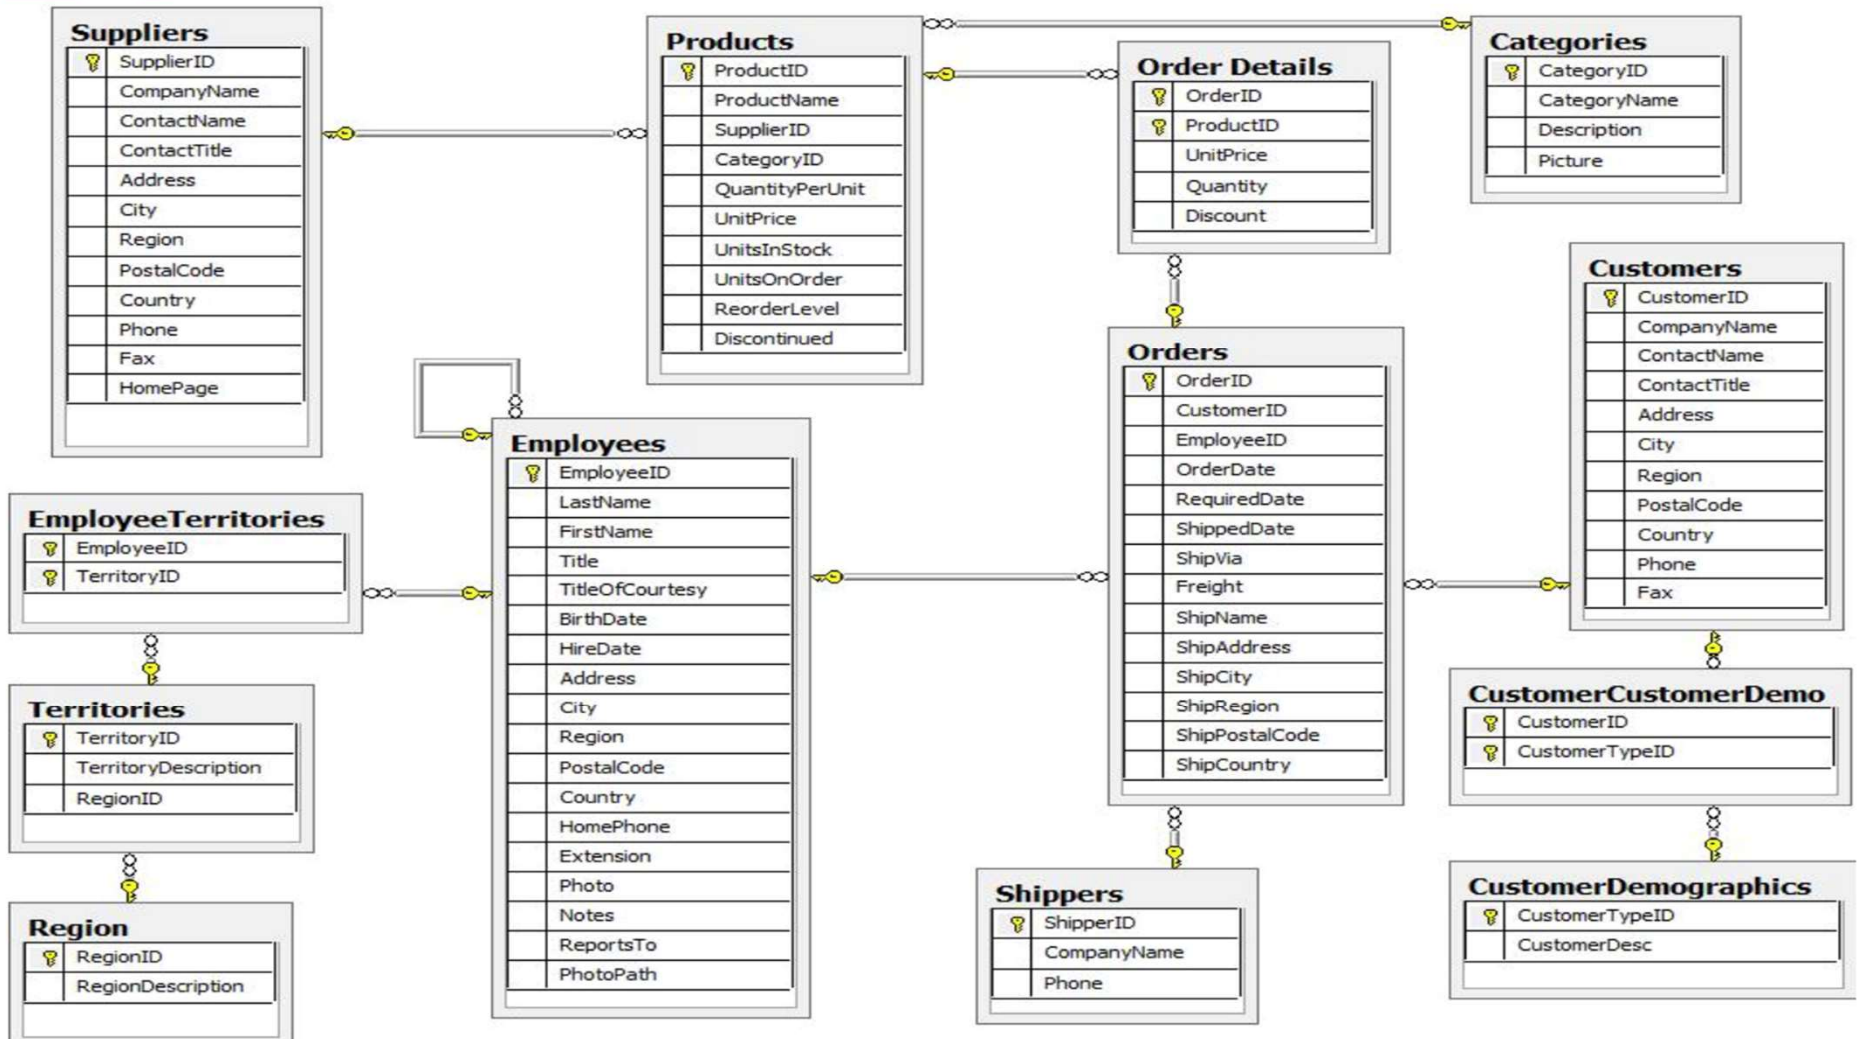


# DB=>Graph Data Model Transformation



Source: [https://neo4j.com/developer/graph-db-vs-rdbms/#\\_from\\_relational\\_to\\_graph\\_databases](https://neo4j.com/developer/graph-db-vs-rdbms/#_from_relational_to_graph_databases)

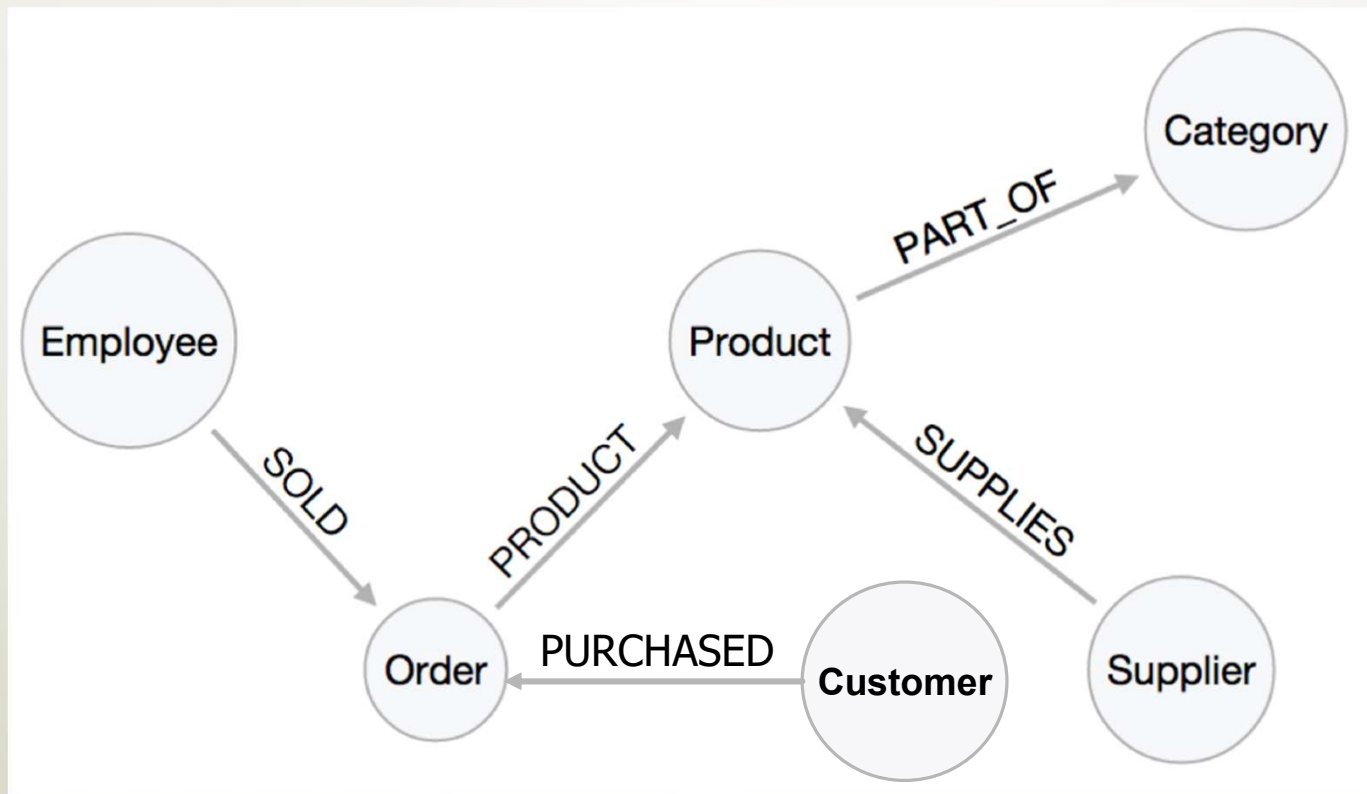
# Northwind Example



Source: <https://neo4j.com/developer/cypher-query-language/>



# Northwind: Graph Model





## Querying Northwind DB: SQL vs. Neo4J

- Select everything from the products table

SQL

```
SELECT p.*  
FROM products as p;
```

Cypher

```
MATCH (p:Product)  
RETURN p;
```



## SQL vs. Neo4J: Projection

- Select Product Name and Price from products table

SQL

```
SELECT p.ProductName, p.UnitPrice  
FROM products as p  
ORDER BY p.UnitPrice DESC  
LIMIT 10;
```

Cypher

```
MATCH (p:Product)  
RETURN p.productName, p.unitPrice  
ORDER BY p.unitPrice DESC  
LIMIT 10;
```



## SQL vs. Neo4J: Filtering

- Select Product Name and Price for “Chocolate”

SQL

```
SELECT p.ProductName, p.UnitPrice  
FROM products AS p  
WHERE p.ProductName = 'Chocolate';
```

Cypher

```
MATCH (p:Product)  
WHERE p.productName = "Chocolate"  
RETURN p.productName, p.unitPrice;
```

OR

```
MATCH (p:Product {productName:"Chocolate"})  
RETURN p.productName, p.unitPrice;
```



## SQL vs. Neo4J: Filtering

- List expensive products that starts with C.

SQL

```
SELECT p.ProductName, p.UnitPrice  
FROM products AS p  
WHERE p.ProductName LIKE 'C%'  
AND p.UnitPrice > 100;
```

Cypher

```
MATCH (p:Product)  
WHERE p.productName STARTS WITH "C"  
AND p.unitPrice > 100  
RETURN p.productName, p.unitPrice;
```





# SQL vs. Neo4J: Joining vs. Traversing

- Who bought Chocolate?

SQL

```
SELECT DISTINCT c.CompanyName  
FROM customers AS c JOIN orders AS o ON (c.CustomerID =  
o.CustomerID)  
JOIN order_details AS od ON (o.OrderID = od.OrderID)  
JOIN products AS p ON (od.ProductID = p.ProductID)  
WHERE p.ProductName = 'Chocolate';
```

Cypher

```
MATCH (p:Product {productName:"Chocolate"})<-  
[:PRODUCT]-(:Order)<-[:PURCHASED]-(c:Customer)  
RETURN distinct c.companyName;
```



## SQL vs. Neo4J: Aggregation

- Find top-selling employees

SQL

```
SELECT e.EmployeeID, count(*) AS Count  
FROM Employee AS e JOIN Order AS o ON  
(o.EmployeeID = e.EmployeeID)  
GROUP BY e.EmployeeID  
ORDER BY Count DESC;
```

Cypher

```
MATCH (:Order)<-[:SOLD]-(e:Employee)  
RETURN e.name, count(*) AS cnt  
ORDER BY cnt DESC;
```



## Summary

- Graph databases store relationships and connections as first-class entities
- Graph databases have good performance when dealing with connected data
- Cypher is a declarative pattern-matching graph query language
- Querying connected data is easier with Cypher



# Neo4j Resources

1. Neo4j Tutorial: <https://www.tutorialspoint.com/neo4j/index.htm>
2. Video Tutorials: [https://neo4j.com/blog/neo4j-video-tutorials/?\\_ga=2.57983406.580712586.1555337212-902296776.1553382068](https://neo4j.com/blog/neo4j-video-tutorials/?_ga=2.57983406.580712586.1555337212-902296776.1553382068)
3. GraphGists are teaching tools which allow you to explore how data in a particular domain would be modeled as a graph and see some example queries of that graph data
  - <https://neo4j.com/graphgists/>
4. Awesome user-defined procedures: <https://github.com/neo4j-contrib/neo4j-apoc-procedures>