



Database Design: ER and UML Diagrams

Abdu Alawini

University of Illinois at Urbana-Champaign

CS411: Database Systems

July 4, 2020

CS411 Goals:


Two Perspectives of DBMS

- USER PERSPECTIVE
 - **how to use a database system?**
 - conceptual data modeling, the relational and other data models, database schema design, relational algebra, SQL and No-SQL query languages.
- SYSTEMS PERSPECTIVE
 - **how to design and implement a database system?**
 - data representation, indexing, query optimization and processing, transaction processing, and concurrency control.
 - NOT COMPLETE: high-level view of implementation; CS511



Overview of Database Design

Today's
lecture

- **Conceptual design:** (ER & UML Models are used for this.)
 - What are the **entities** and **relationships** we need?
- **Logical design:** 
 - Transform ER design to Relational Schema
- **Schema Refinement:** (Normalization)
 - Check relational schema for redundancies and related anomalies.
- **Physical Database Design and Tuning:**
 - Consider typical workloads; (sometimes) modify the database design; select file types and indexes.

We'll do
this next lecture



Entity-Relationship Model is a different model than the Relational Model

- **Relational model** has:
 - **tables** (relations) with attributes, keys, foreign keys, domain definitions for attributes
- **Entity-Relationship model** has:
 - **Entities and entity sets** with attributes, keys, and domain definitions for attributes
 - **Relationships among entities and relationship sets** with uniqueness or cardinality constraints



Entity Relationship Model Unified Modeling Language

ER Model

- Proposed by Peter Chen in 1976
- Gives us a language to specify
 - What information the database must hold
 - How the bits of information relate to one another

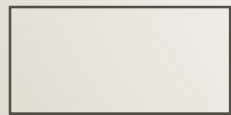
UML Model

- UML is a standard language for designing software systems
 - also used for DB design
- created by the Object Management Group (OMG)
- UML 1.0 specification draft was proposed to the OMG in early 1997.

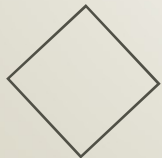


Entity-Relationship Diagram (original syntax)

Legend:



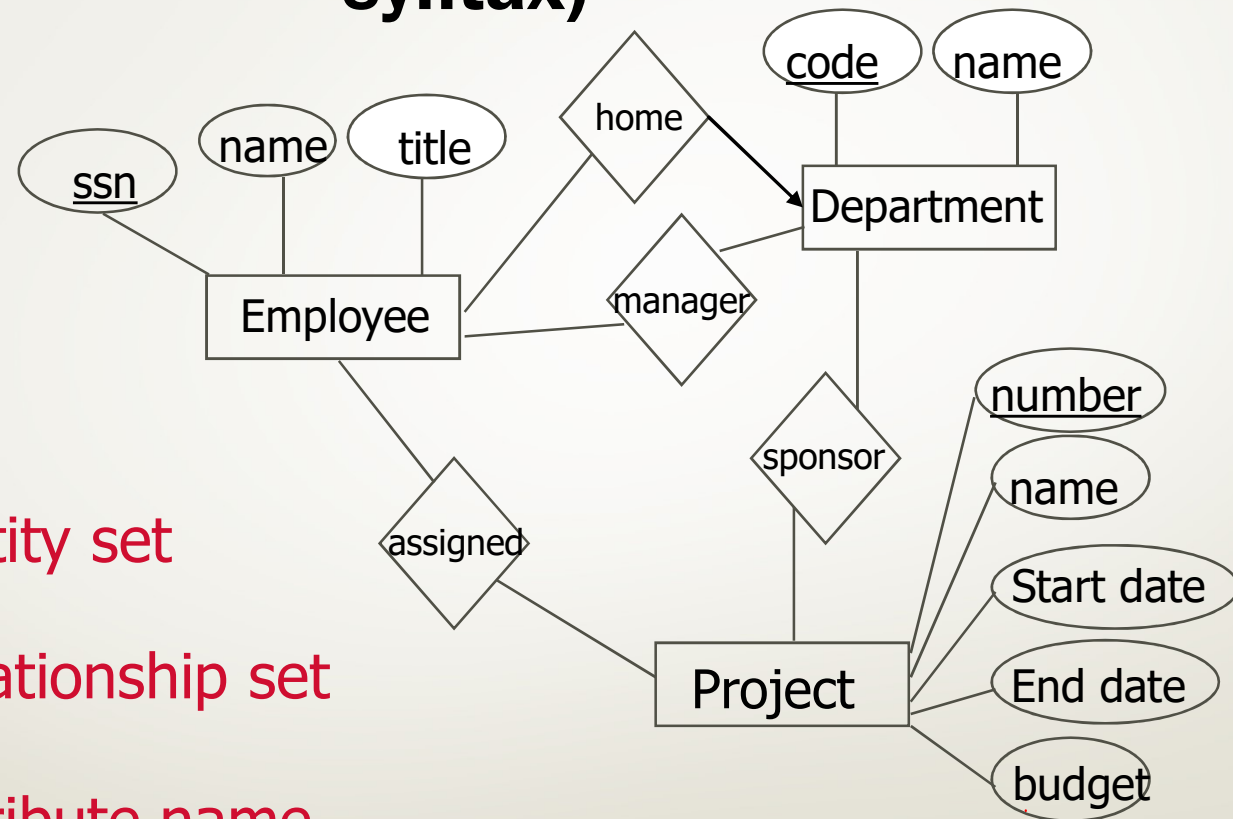
Entity set



Relationship set



Attribute name



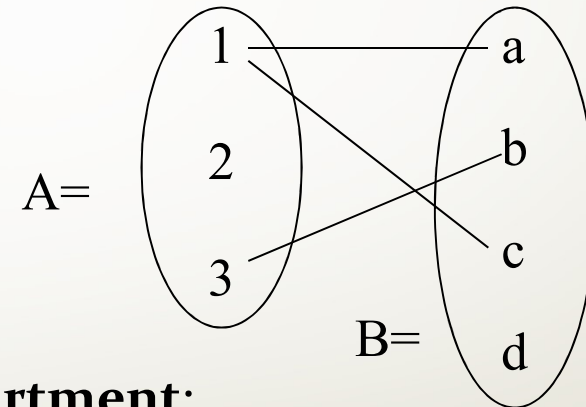


Definitions

- Entity: Real-world object distinguishable from other objects.
An entity is described using a set of *attributes*.
- Entity Set: A collection of similar entities. E.g., all employees.
(often referred to as just entity, which blurs the distinction between type and collection)
- Relationship: Association among 2 or more entities. E.g., Kristin's *home department* is Research & Development.
- Relationship Set: Collection of similar relationships. E.g., Home
(often referred to as just relationship)

Relationships

- Formal definition:
 - if A, B are sets, then a relation R is a subset of $A \times B$
- $A = \{1, 2, 3\}$, $B = \{a, b, c, d\}$,
 $R = \{(1, a), (1, c), (3, b)\}$



Same story w/ entity sets

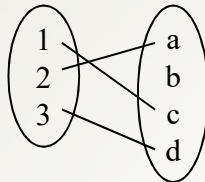
sponsor is a subset of **Project** x **Department**:



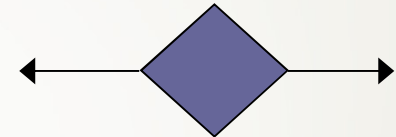


Multiplicity of ER Relationships

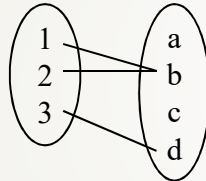
- one-one:



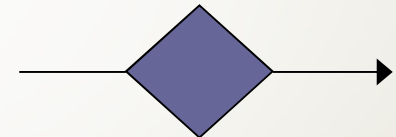
One on LHS/RHS
connected to at most
one on RHS/LHS



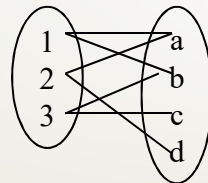
- many-one:



One on LHS
connected to at most
one on RHS



- many-many:



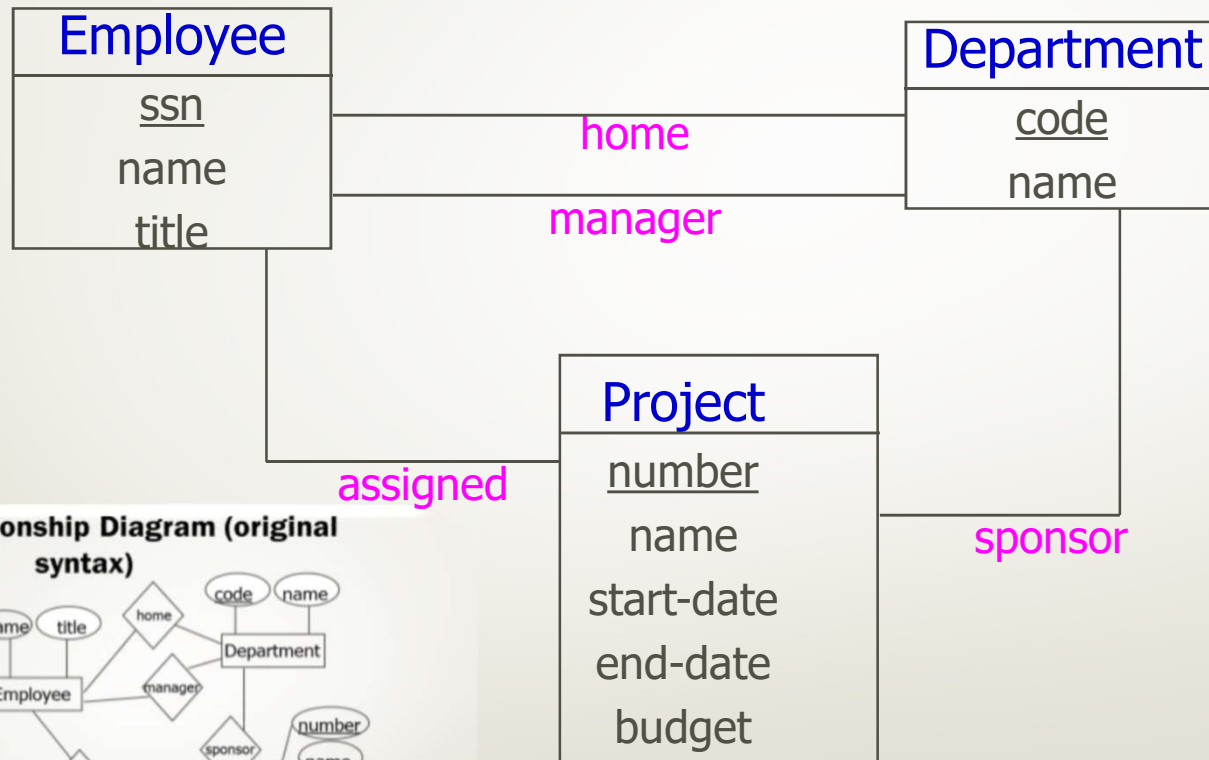
No constraints



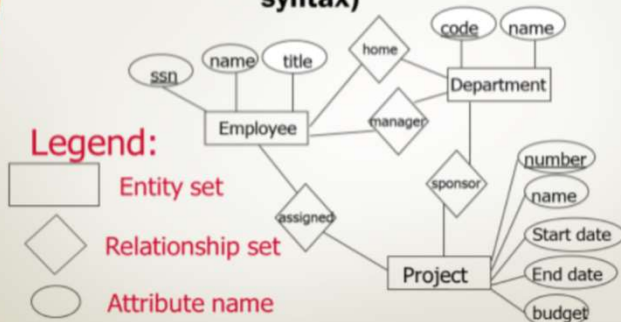
- ❑ Multiplicity can be shown with arrows
- ❑ Arrow = at most 1



UML version of the same E-R Diagram

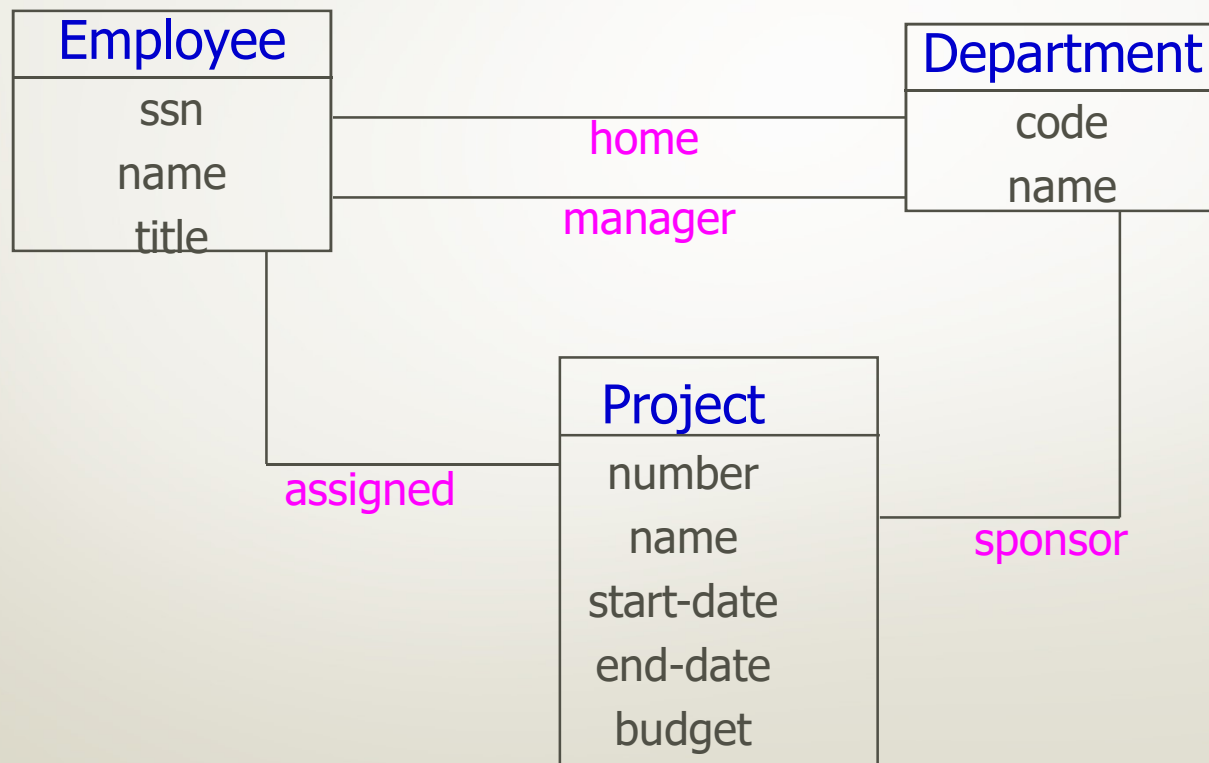


Entity-Relationship Diagram (original syntax)



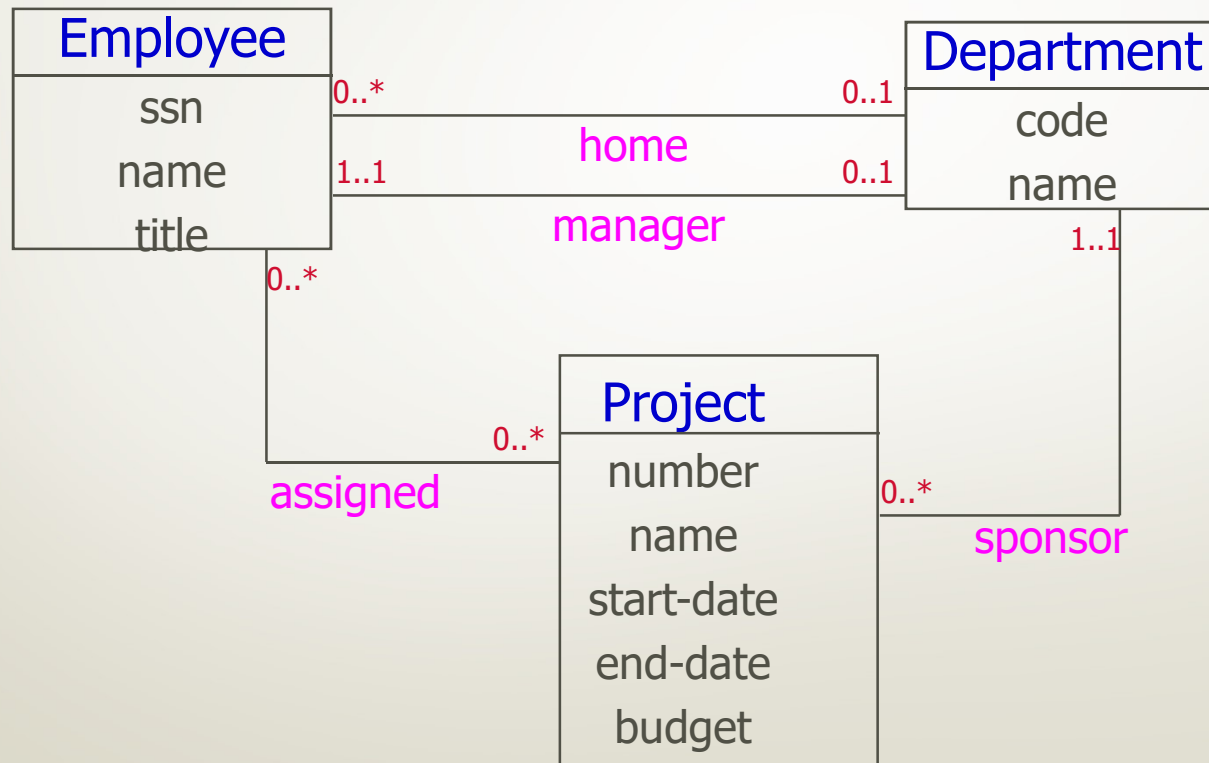


Cardinality Constraints on Relationship sets: How many entities can participate?



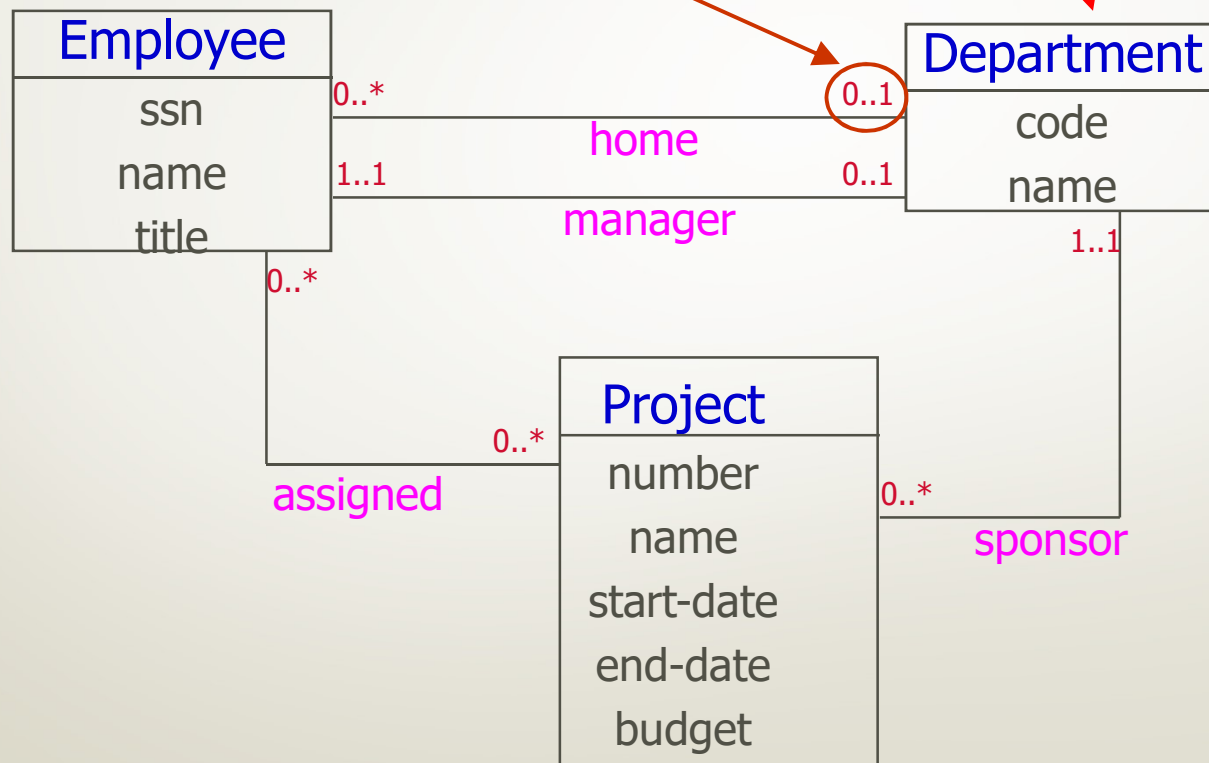


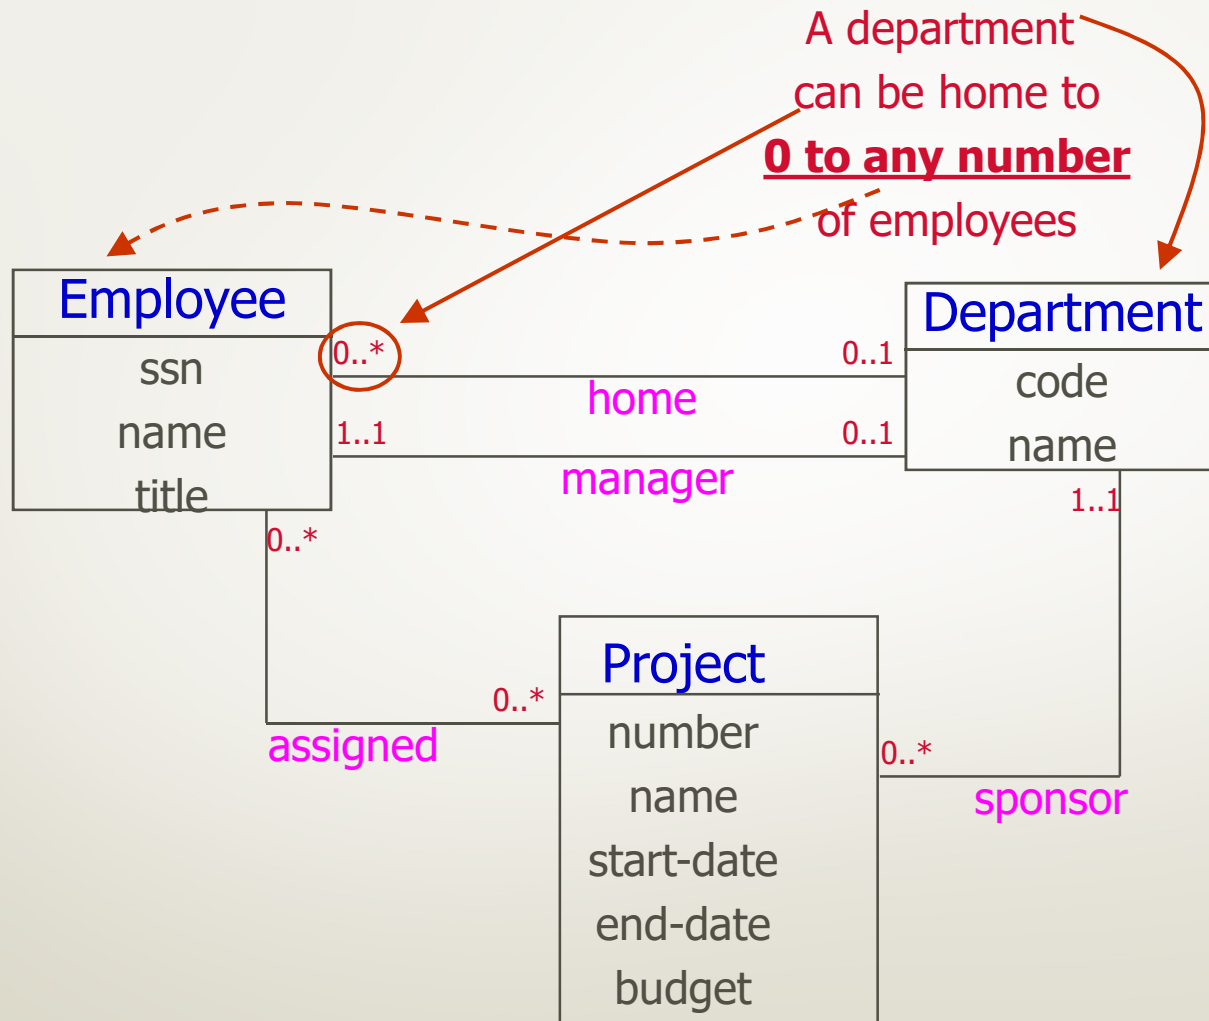
Cardinality Constraints on Relationship sets: How many entities can participate?





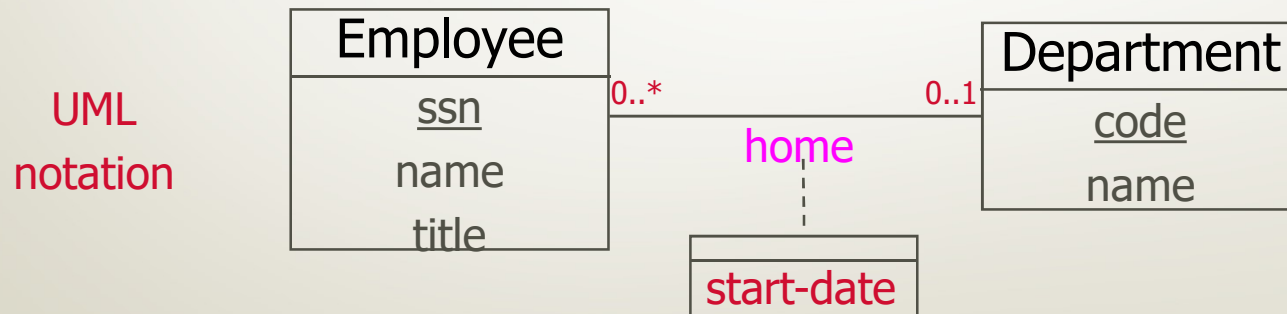
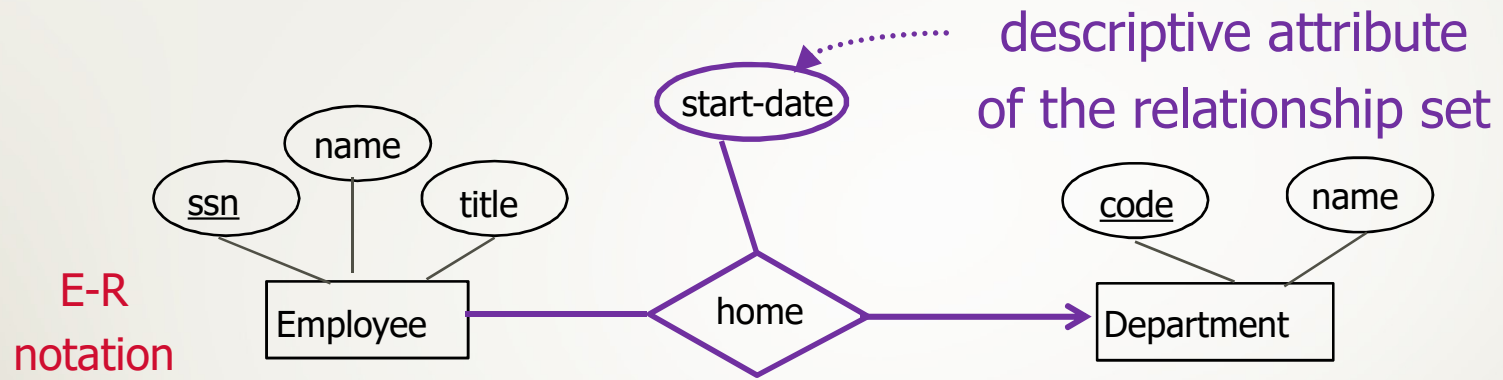
An employee can have **0 or 1** home departments







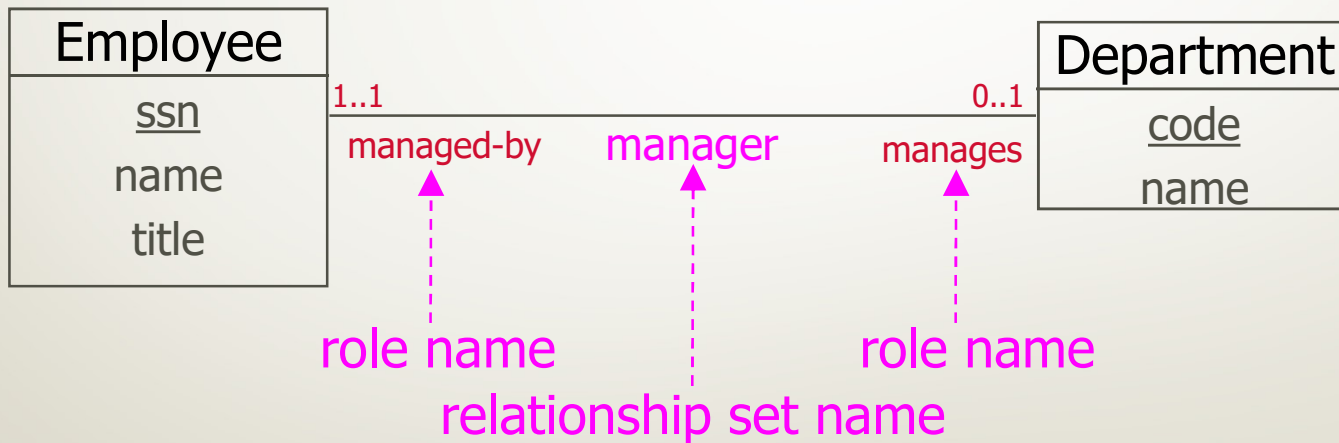
Relationship sets can have attributes





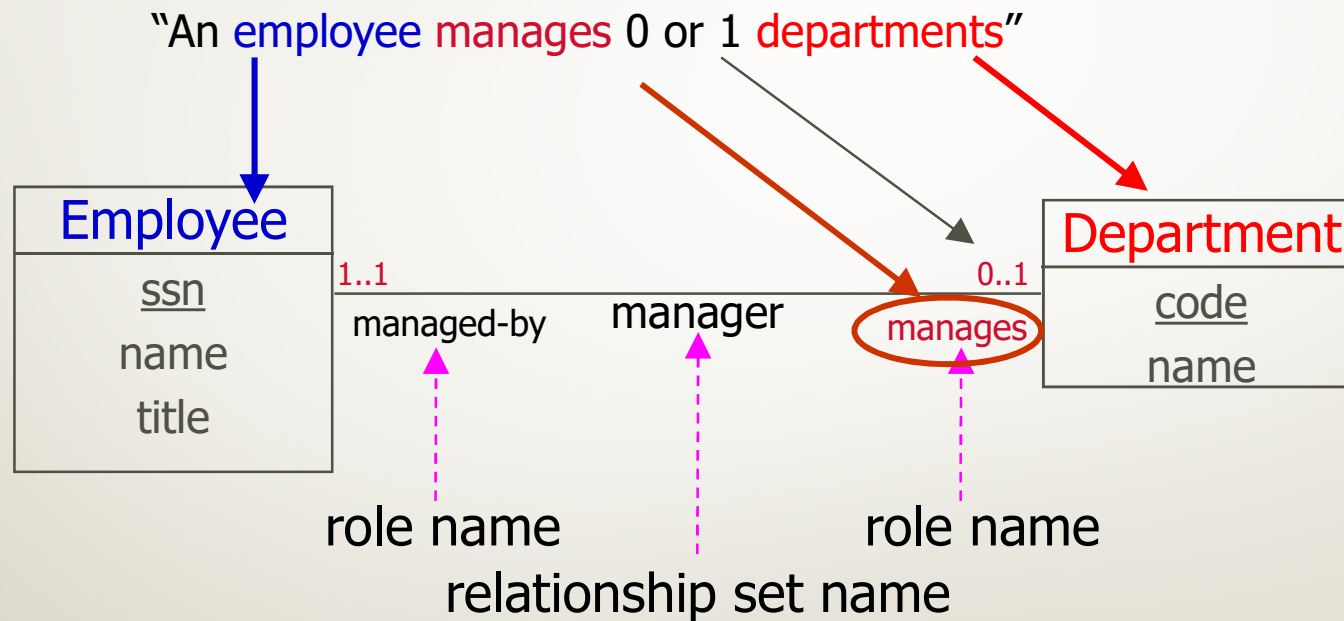
Relationship sets can have **role** names

(in addition to the name of the relationship set)





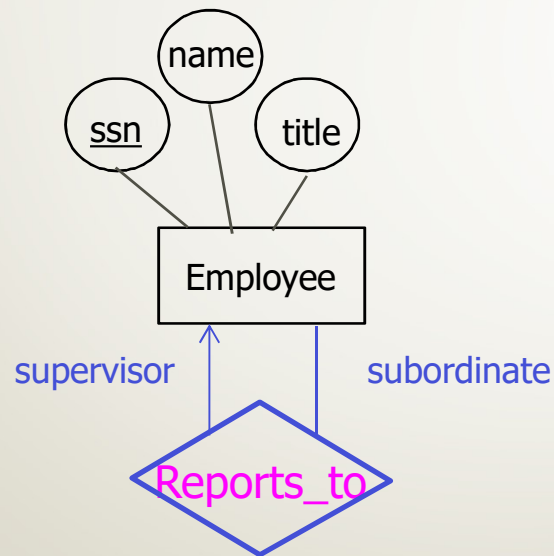
Example: reading **role** names



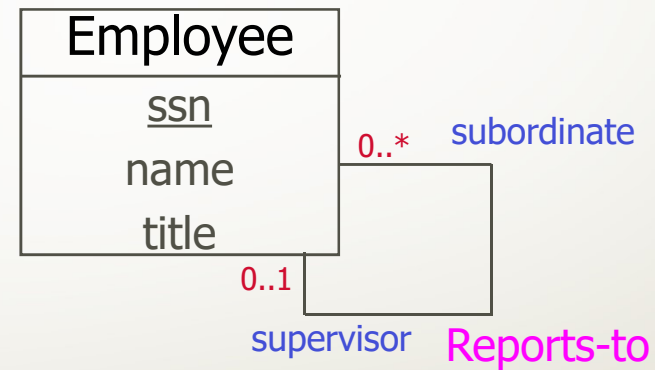


Same entity sets can participate in different “roles” for the same relationship set

E-R
notation



UML
notation





Constraints in ER diagram

- Recall that a constraint is an assertion about the database that must be true at all times
- Part of the database schema = structure
(so it must be part of the ER diagram)
- Very important in database design



Modeling Constraints

Finding constraints is part of the modeling process.
Commonly used constraints:

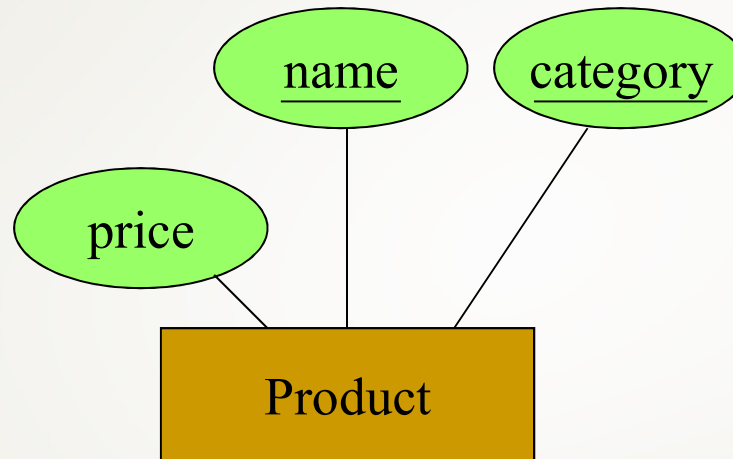
- ❑ **Keys:** attributes that identify entities in an entity set
e.g., social security number uniquely identifies a person.
- ❑ **Referential integrity constraints:** relationship-based constraints
e.g., if you work for a company, it must exist in the database.
- ❑ **Domain constraints:** peoples' ages are between 0 and 150.
- ❑ **General constraints:** all others (at most 50 students enroll in a class)



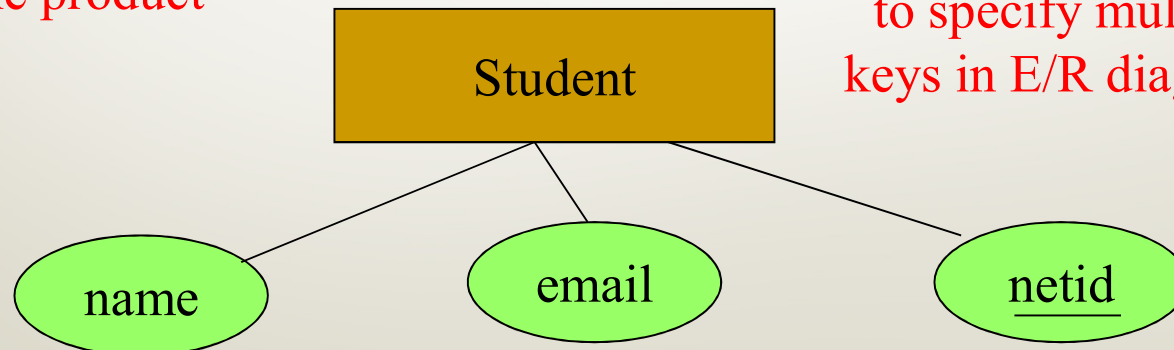
Keys in E/R Diagrams

Underline:

This means “name”
and “category”
together uniquely
determine product



No formal way
to specify multiple
keys in E/R diagrams

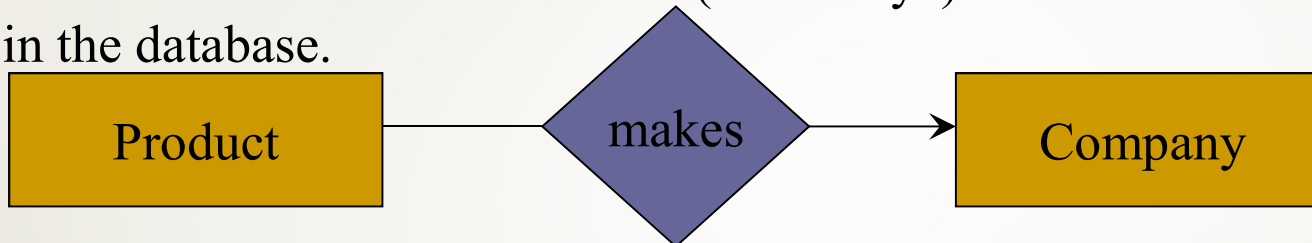




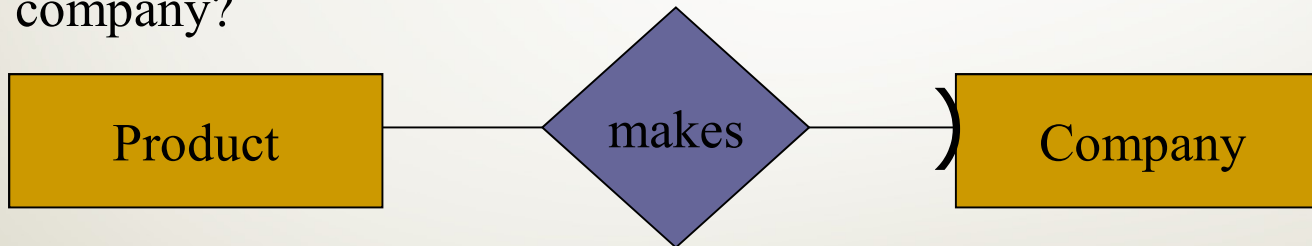
Referential Integrity Constraints

Recall: the arrow meant “at most one”.

Each Product must be related to (“made by”) at most one Company in the database.



Wouldn't it be weird if a product was not associated with any company?



This says “exactly one”.

Each Product must be related to (“made by”) exactly one Company in the database.

Arrow = at most 1

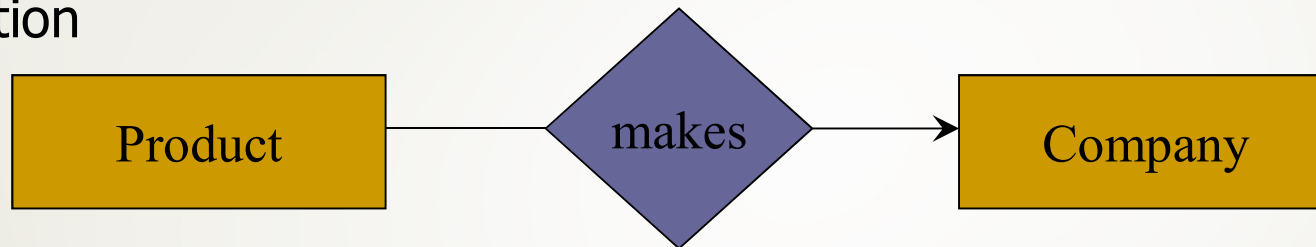
Semicircle = exactly 1



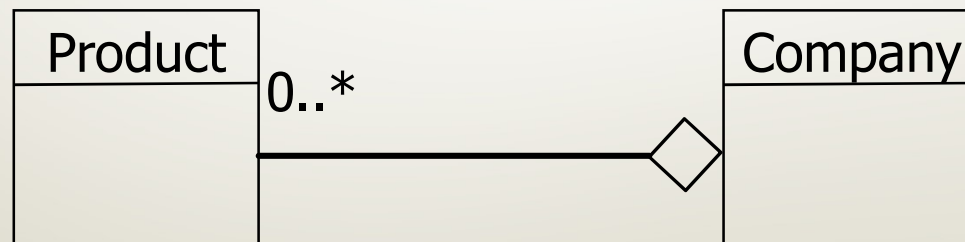
Referential Integrity Constraints

UML: Aggregation

ER
Notation



UML
Notation

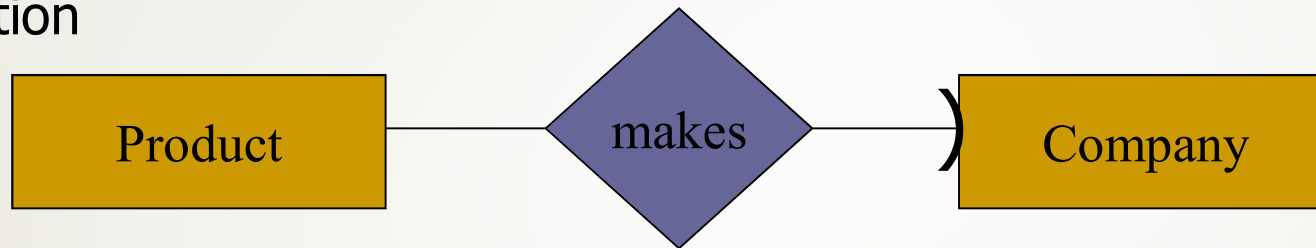




Referential Integrity Constraints

UML: Composition

ER
Notation



UML
Notation

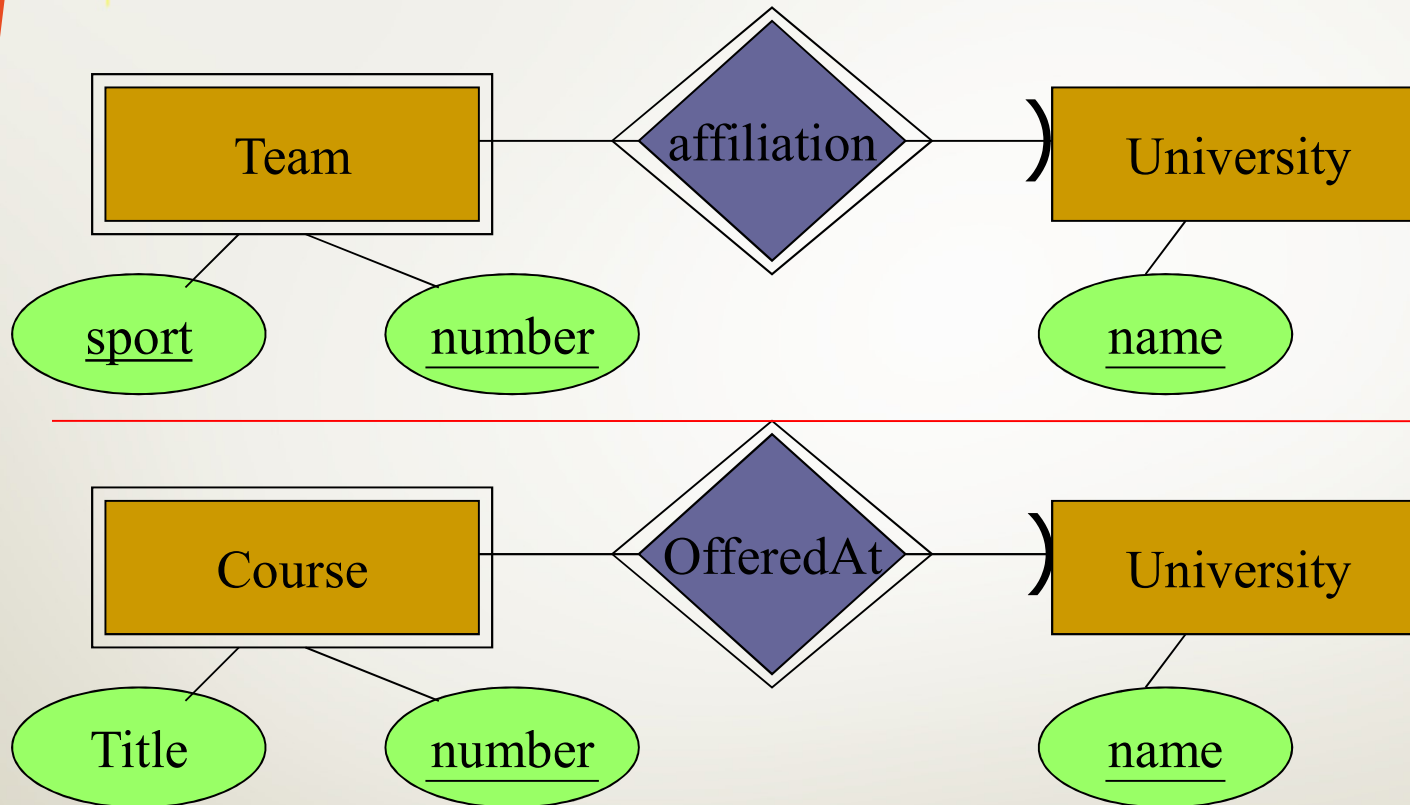




Weak Entity Sets

- Occasionally, entities of an entity set need “help” to identify them uniquely.
- Entity set E is *weak* if in order to identify entities of E uniquely, we need to follow one or more many-one relationships from E and include the key of the related entity sets.
- Note: not an is-a relationship because E is not a “subclass” of F :
Univ and Team

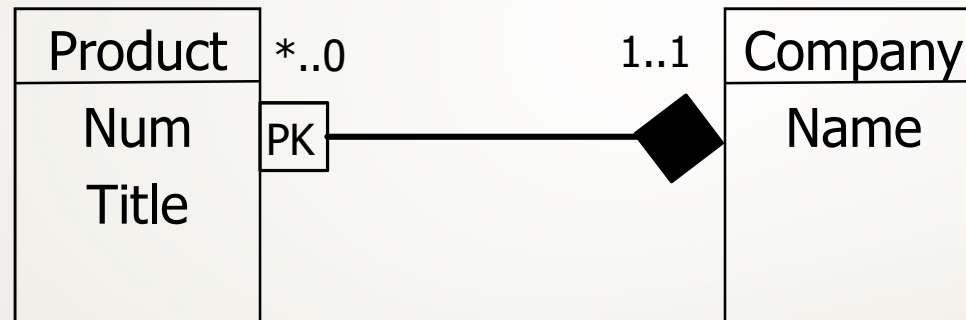
Notations for weak entity set



- “University” is a “supporting entity set” for “Team”.
- “Affiliation” is a “supporting relationship”.



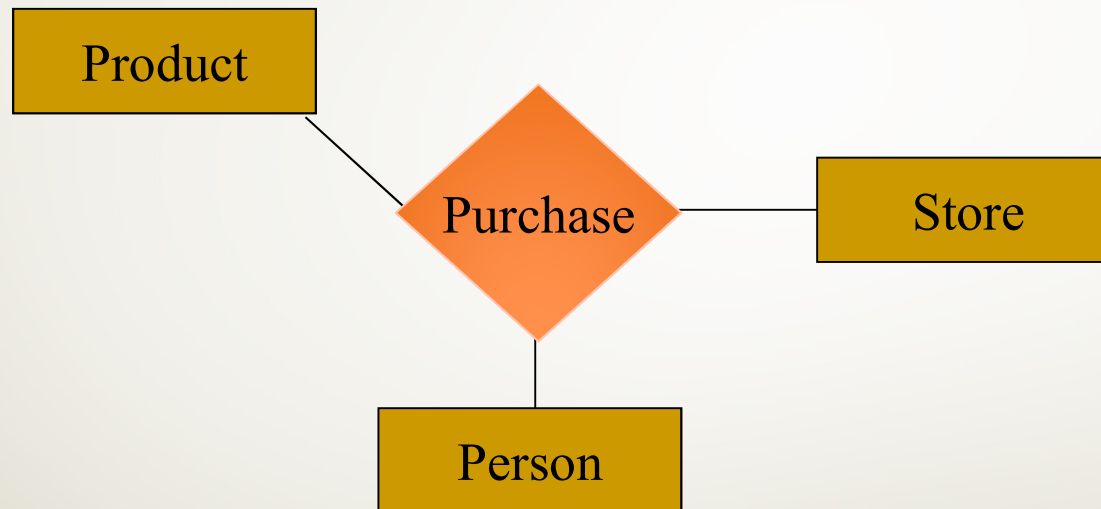
Weak entity set in UML





Multiway Relationships

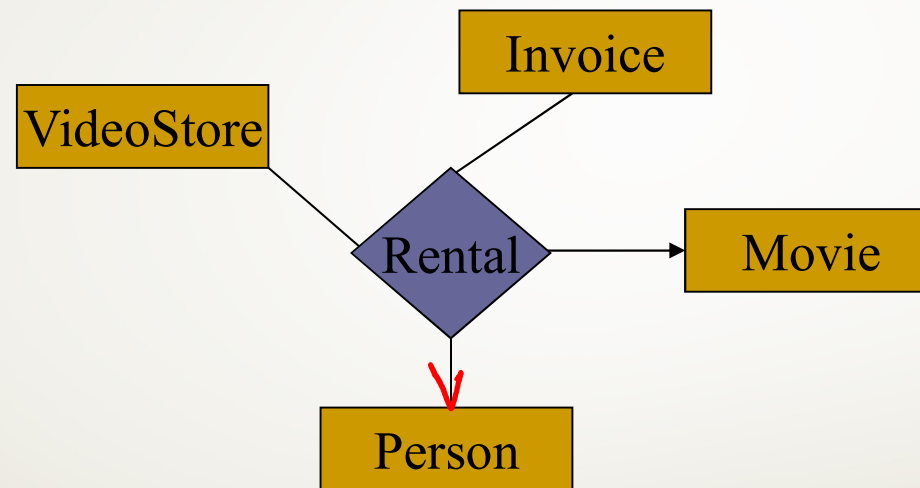
How do we model a purchase relationship between buyers, products and stores?



- 10 Can still model as a mathematical set (how ?)
 - Yes: As a subset of product x store x person

Arrows in Multiway Relationships

Q: what does the arrow mean ?



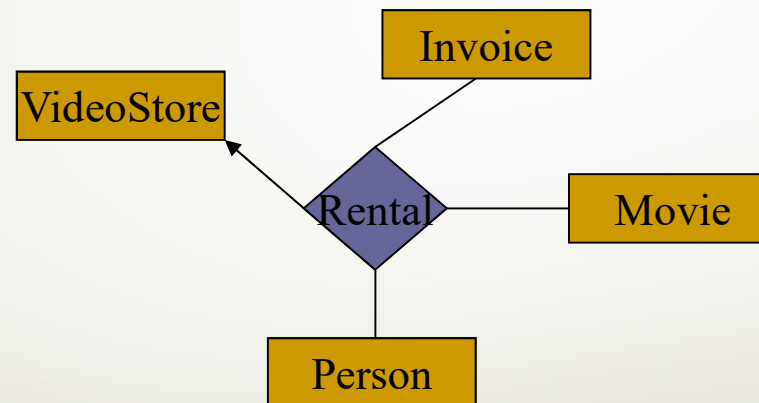
A: “At most one”. That is, a specific combination of videostore, invoice and person can correspond to at most one movie.

Q: What if I had an arrow into Person?

Arrows in Multiway Relationships

Q: how do I say: “invoice determines store” ?

A: no good way; best approximation:



Q: Why is this bad ?

A: We aren't clarifying that the store is a function of the invoice only

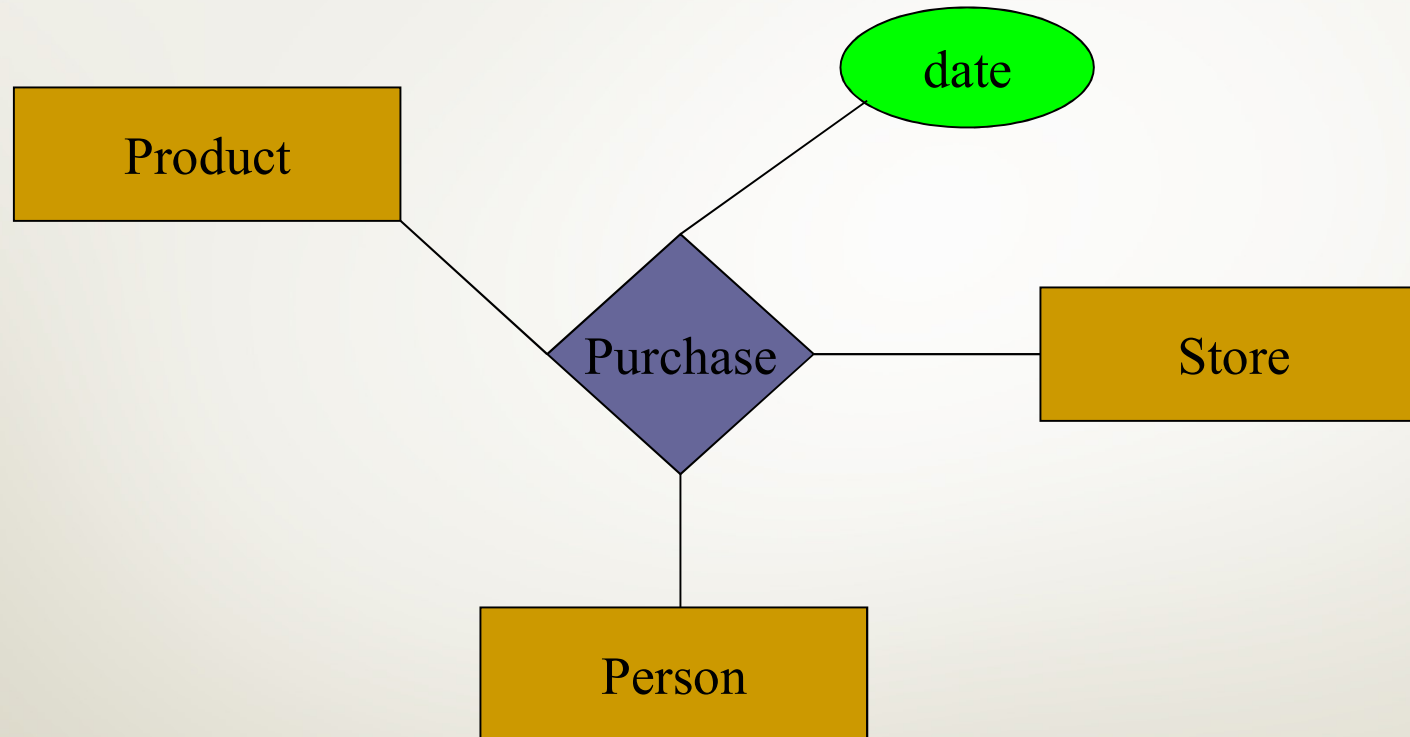


Some ER Modeling Tools Require 2-way Relationships

Do we need multi-way relationships or
do 2-way (binary) relationships suffice?

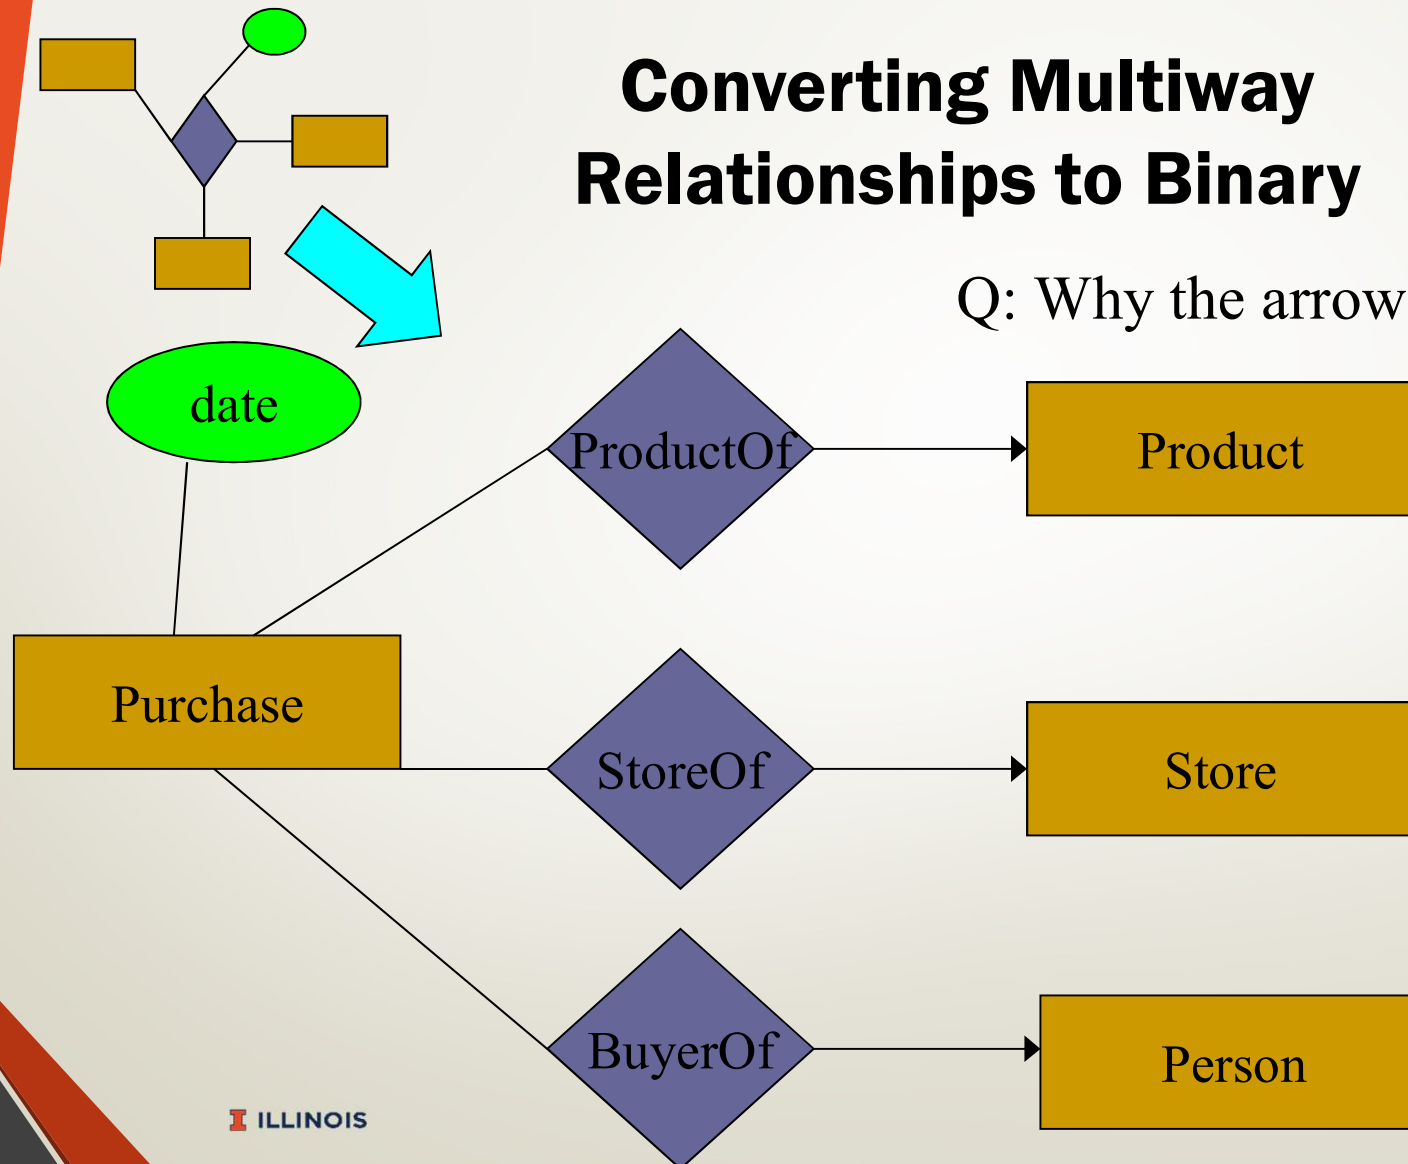


How would you convert this into binary?



Converting Multiway Relationships to Binary

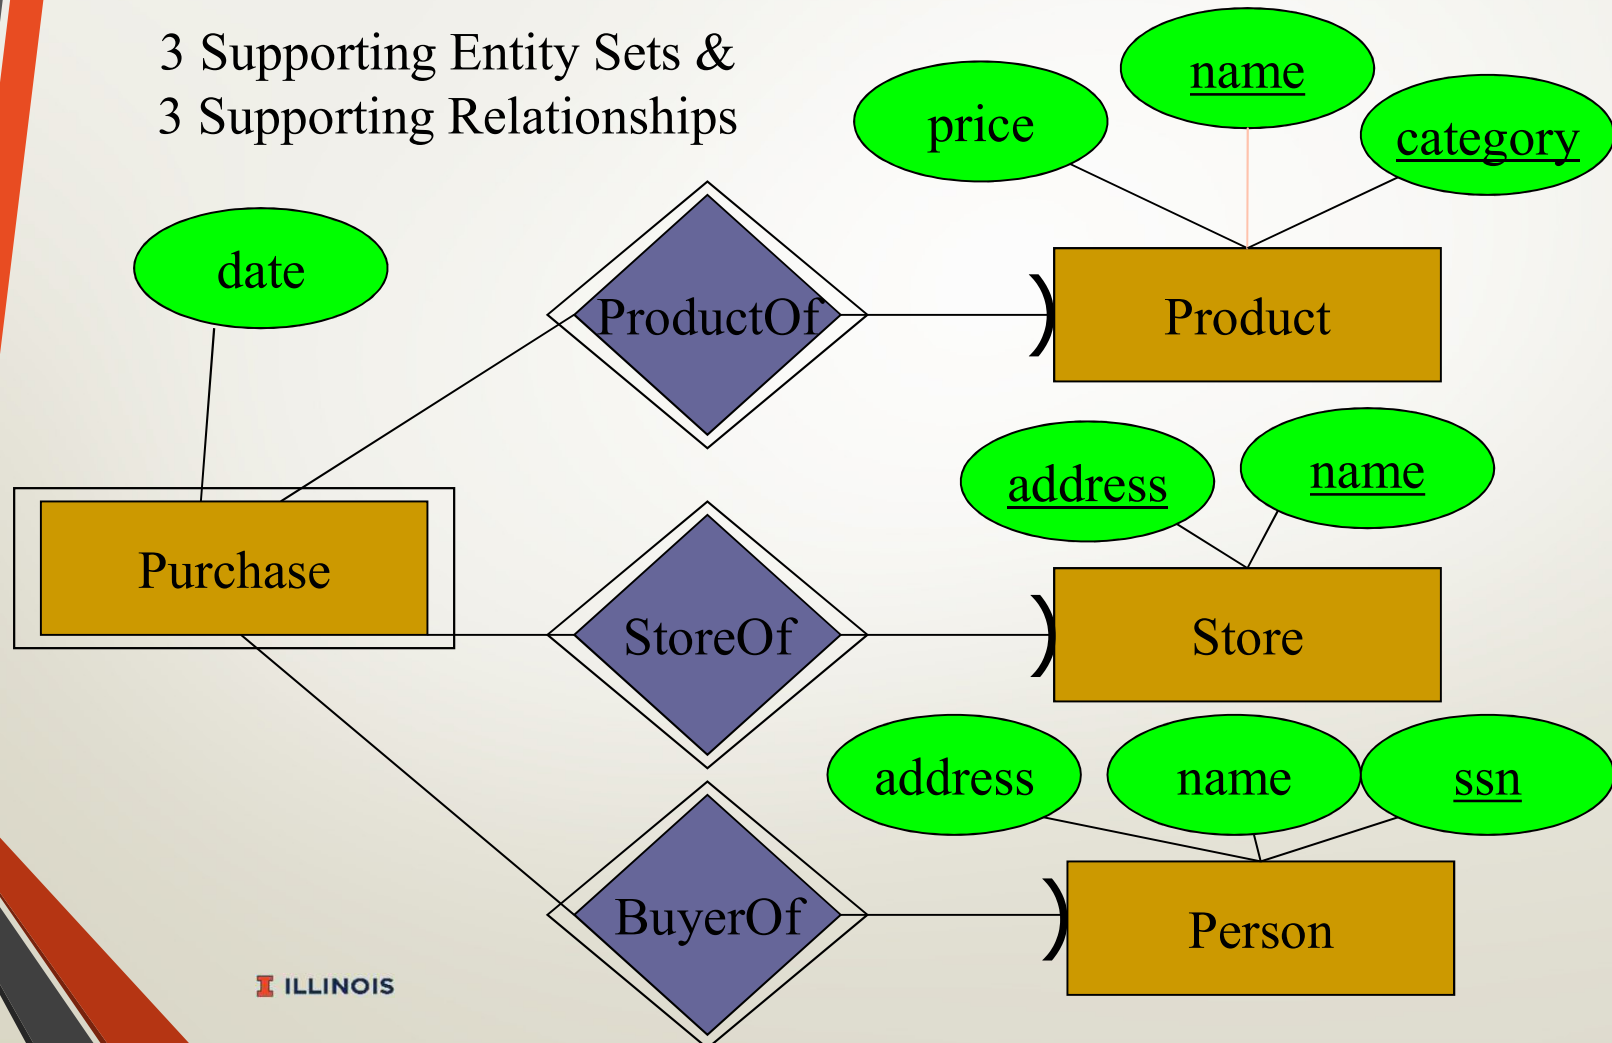
Q: Why the arrows?





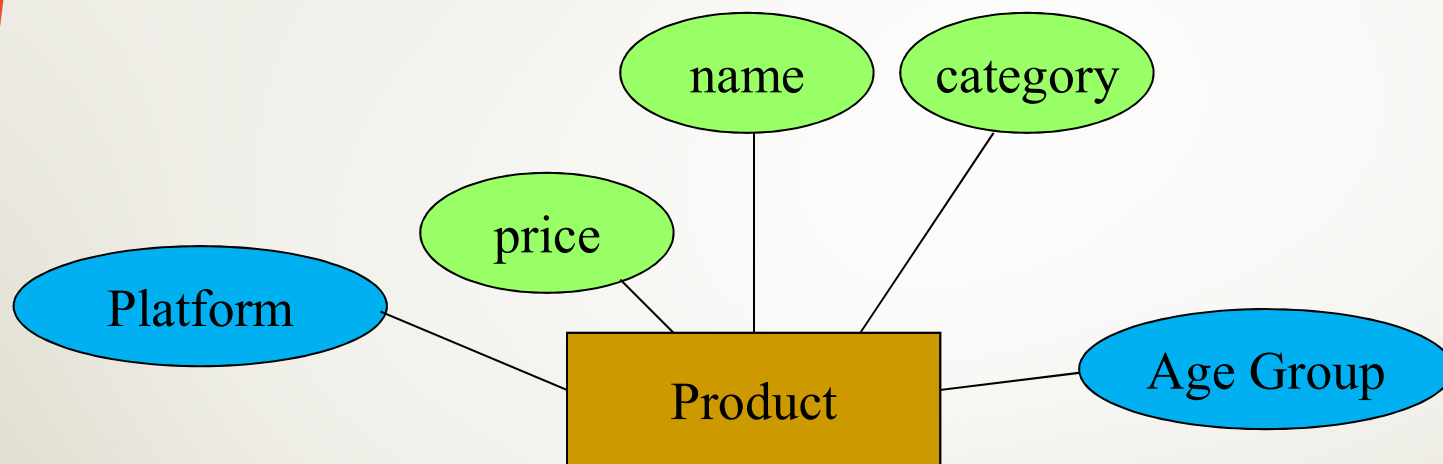
Another scenario where weak e.s. arises

3 Supporting Entity Sets &
3 Supporting Relationships

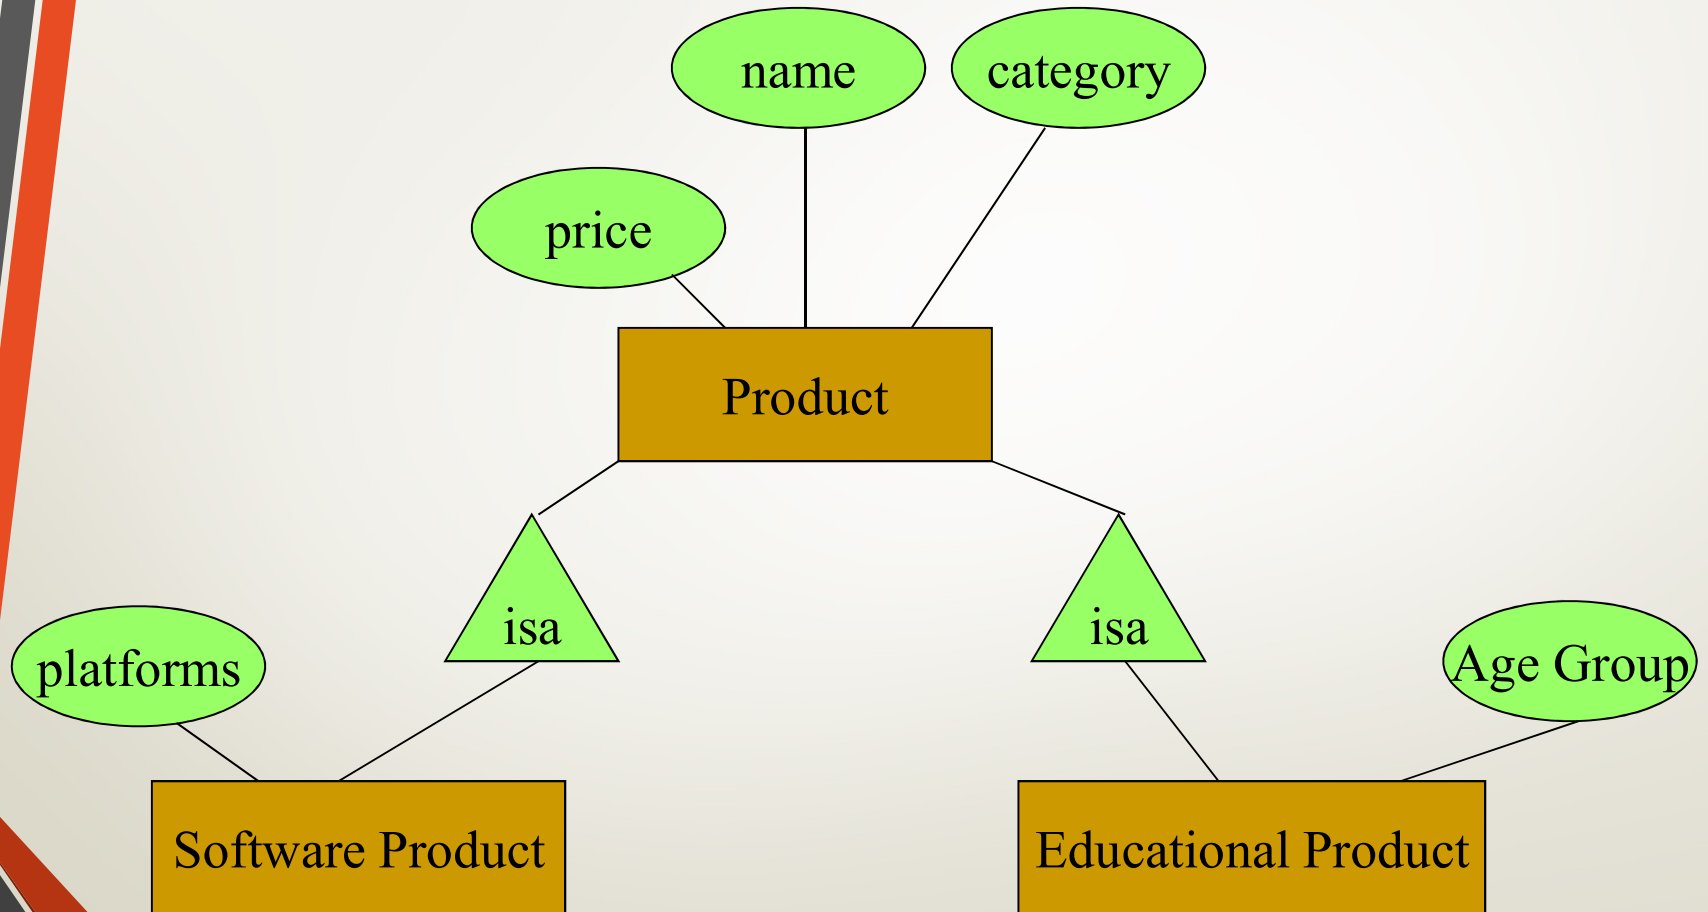




Subclasses in ER Diagrams

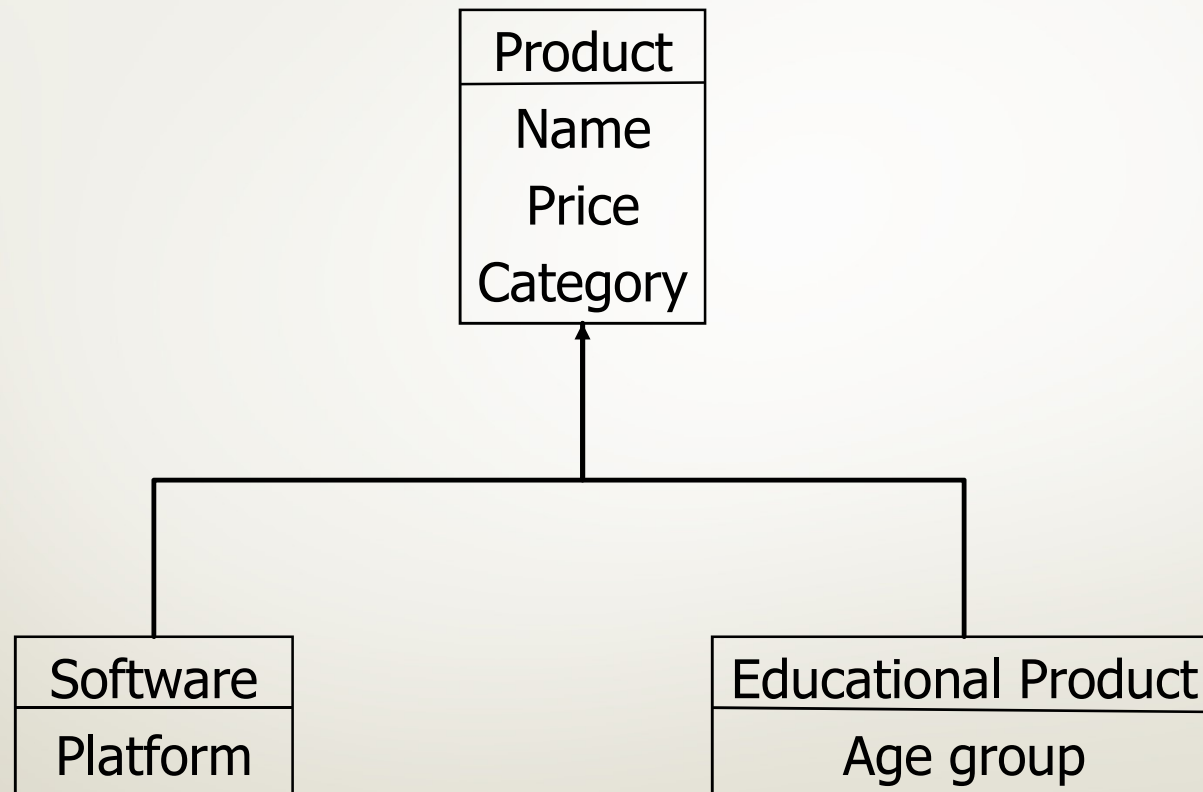


Subclasses in ER Diagrams





Subclasses in UML





Subclasses

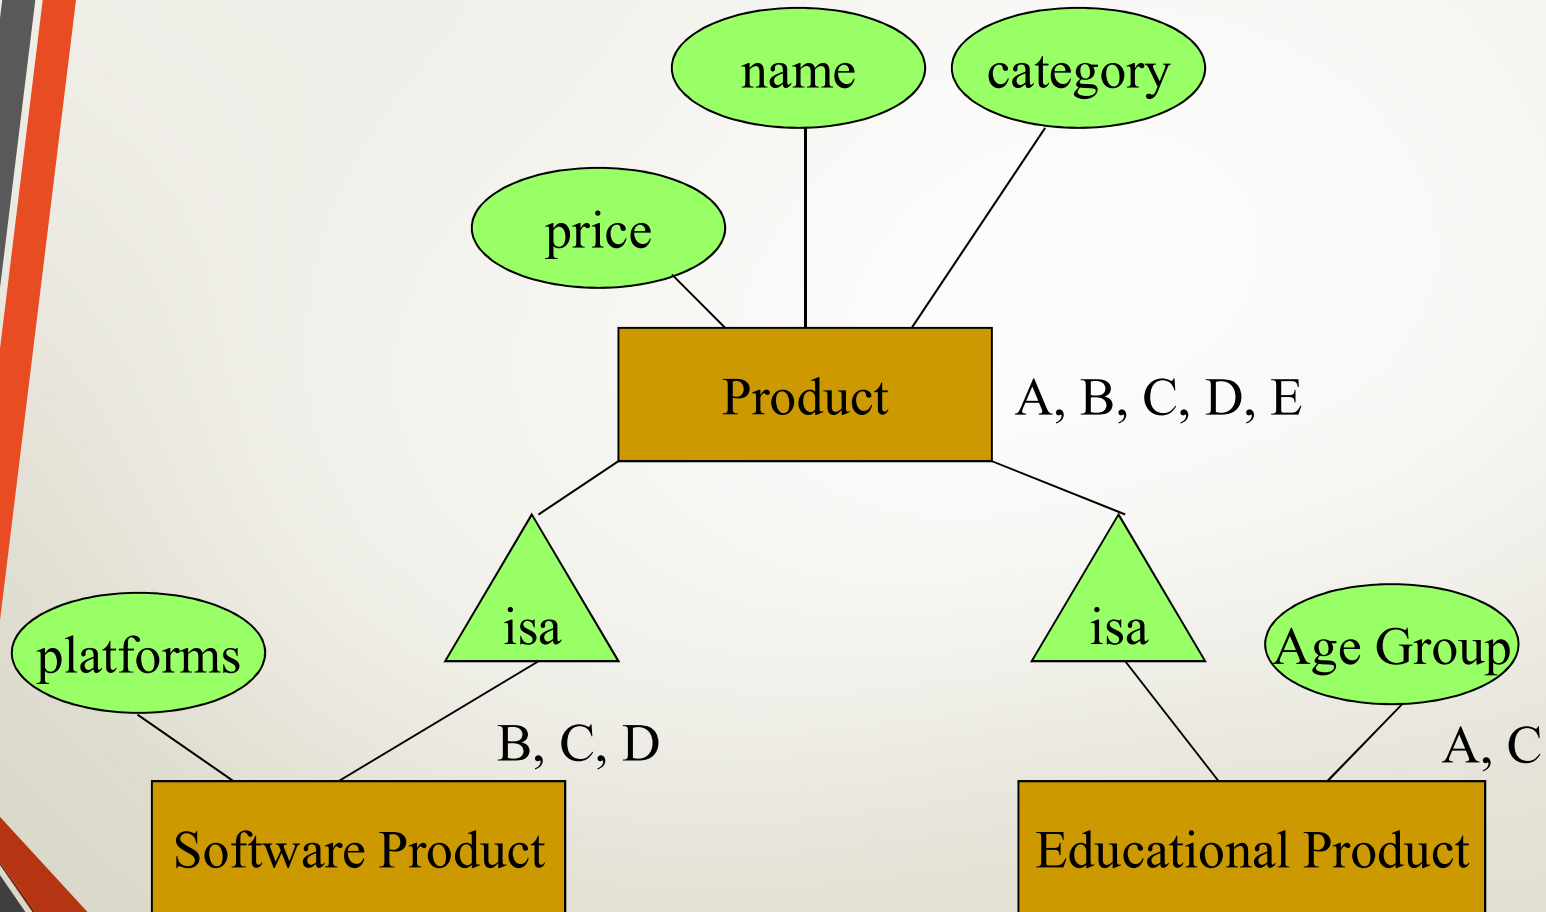
- “Isa” triangles indicate the subclass relationship.
 - Point to the superclass.
- Subclasses form a tree.
 - I.e., no “multiple inheritance”.
- Why subclasses?
 - Unnecessary to add redundant properties to the root entity set that don’t apply to many of the entities



ER Vs. Object Oriented Subclasses

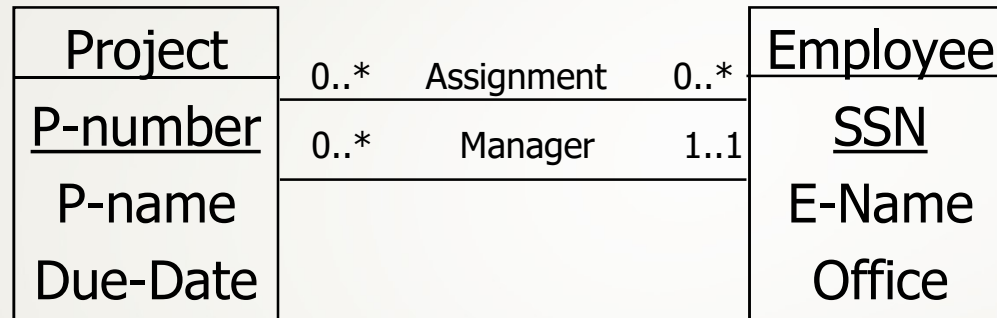
- In the object-oriented world, objects are in one class only.
 - Subclasses inherit properties from superclasses.
- In contrast, E/R entities have components in all subclasses to which they belong.
 - Matters when we convert to relations.

Subclasses in ER Diagrams





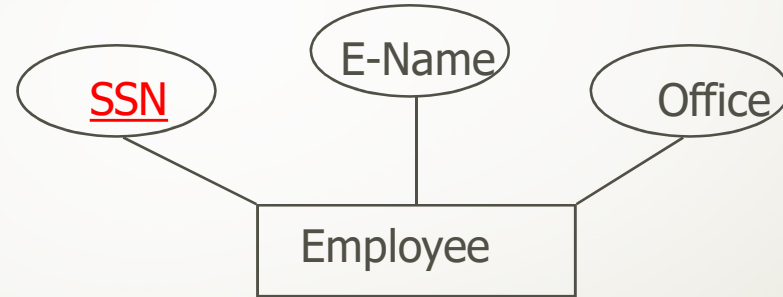
Converting ER to Relational Schema





1. Translate each entity set into a table, with keys.

- **Entity set:**
 - can be represented as a table in the relational model
 - has a **key** ... which becomes a key for the table

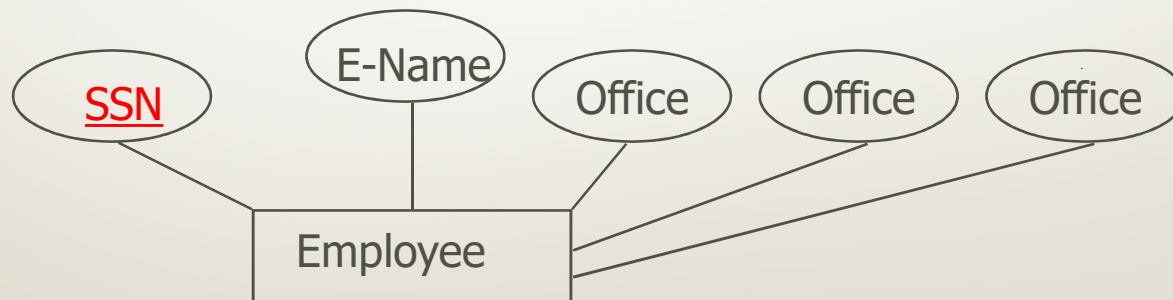


```
CREATE TABLE Employee
(SSN CHAR(11) NOT NULL,
E-Name CHAR(20),
Office INTEGER,
PRIMARY KEY (SSN))
```



Multi-valued Attribute

Didn't see this case when discussing ER diagrams
One or more values of same attribute for an entity



I

2. Create a table for the multi-valued attribute.

Most relational DBMSs do not allow multi-valued attributes.

How many offices can one employee have?

Just one

Project(P-number, P-name, Due-Date)
Employee(SSN, E-Name, Office)

VS.

More than one

Project(P-number, P-name, Due-Date)
Employee(SSN, E-Name)
Office-Assignment(SSN, Office)



Sample Data

Project(P-number, P-name, Due-Date)

Employee(SSN, E-Name, **Office**)

*Just
one*

12 Smith O-105

15 Wei O-110

Project(P-number, P-name, Due-Date)

Employee(SSN, E-Name)

12 Smith

15 Wei

Office-Assignment(SSN, Office)

12 O-105

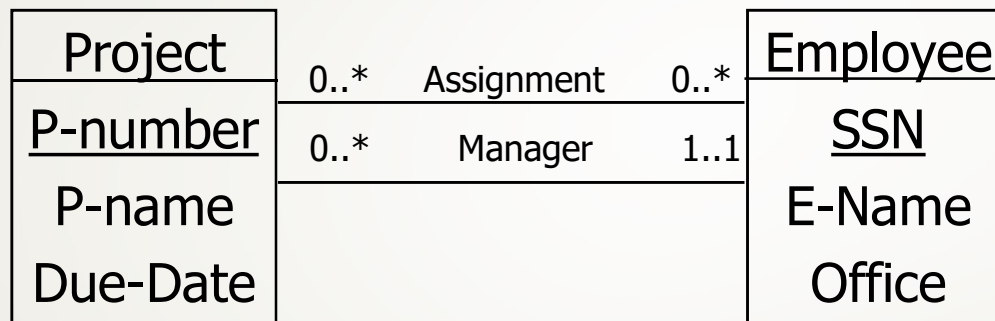
12 O-106

15 O-110

*More
than
one*



3. Translate each **many-to-many** relationship set into a table



What are the attributes and what is the key for Assignment?

Project(P-number, P-name, Due-Date)

Employee(SSN, E-Name, Office)



3. Translate each **many-to-many** relationship set into a table

Project	0..*	Assignment	0..*	Employee
<u>P-number</u>	0..*	Manager	1..1	<u>SSN</u>
P-name				E-Name
Due-Date				Office

Answer: Assignment(P-Number, SSN)

P-Number is a foreign key for Project

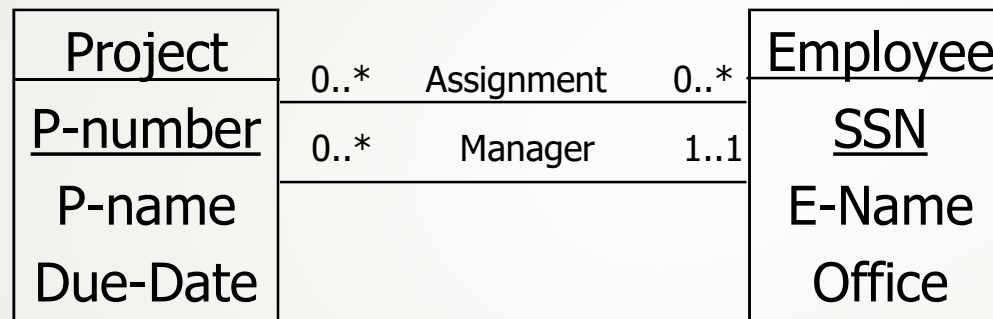
SSN is a foreign key for Employee

Project(P-Number, P-Due-Date)

Employee(SSN, E-Name, Office)



What should we do with each **one-to-many** relationship set?



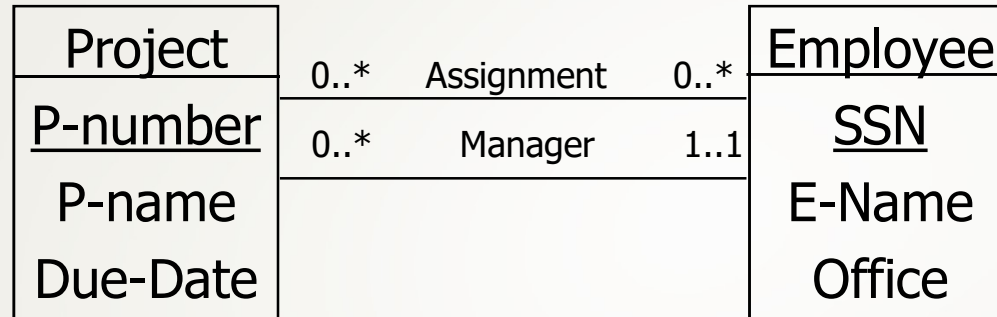
Manager (?)

Project(P-number, P-name, Due-Date)

Employee(SSN, E-Name, Office)



4. Create a foreign key for a **1-to-many** relationship set.



Project(P-number, P-name, Due-Date, **MgrSSN**)
Employee(SSN, E-Name, Office)

MgrSSN is a foreign key (referencing the Employee relation)

value of Manager must match an SSN



4. Create a foreign key for a **1-to-many** relationship set.

Project	0..*	Assignment	0..*	Employee
<u>P-number</u>	0..*	Manager	1..1	<u>SSN</u>
P-name				E-Name
Due-Date				Office

Project(P-number, P-name, Due-Date, MgrSSN)

Employee(SSN, E-Name, Office)

VS.

Project(P-number, P-name, Due-Date)

Employee(SSN, E-Name, Office)

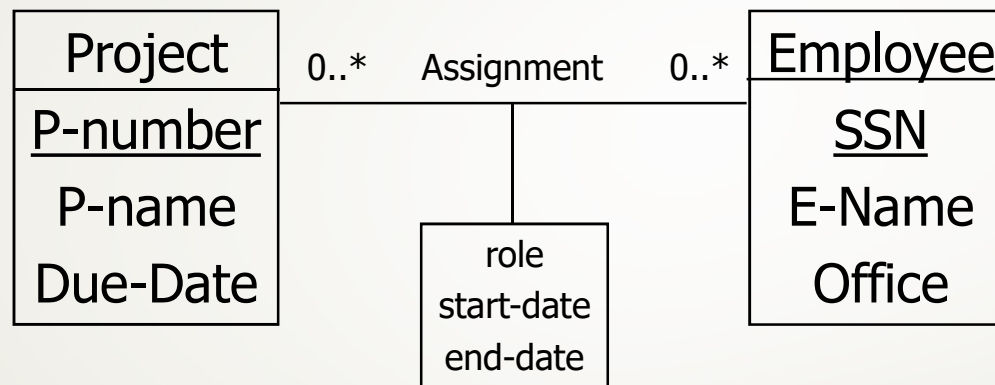
Manager(P-number, SSN)

What are the tradeoffs between these two?

Note:
P-number
is the key
for Manager



What do we do when a **many-to-many** relationship set has an attribute?



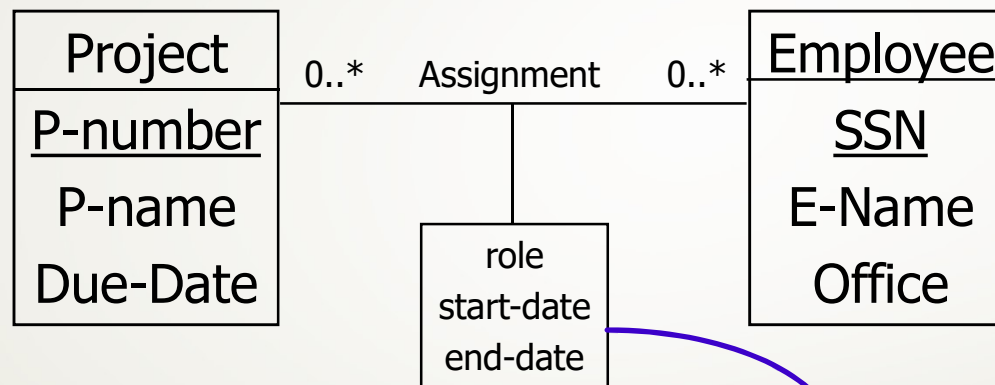
Assignment(P-number, SSN)

Project(P-number, P-name, Due-Date)

Employee(SSN, E-Name, Office)



What do we do when a **many-to-many** relationship set has an attribute?



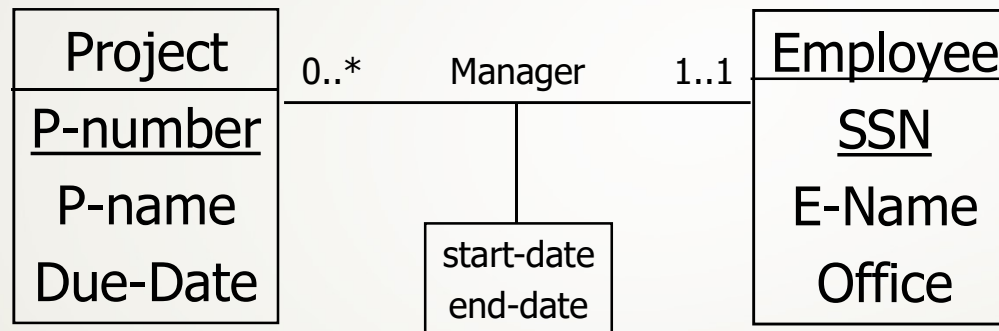
Assignment(P-number, SSN, role, start-date, end-date)

Project(P-number, P-name, Due-Date)

Employee(SSN, E-Name, Office)



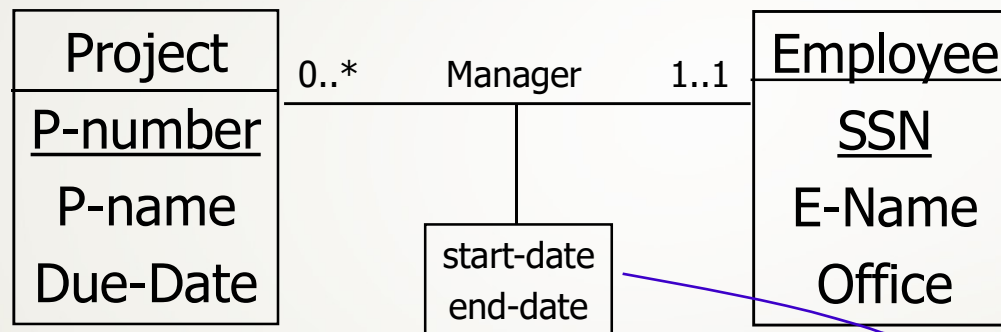
What do we do when a **1-to-many** relationship set has an attribute?



Project(P-number, P-name, Due-Date, MgrSSN)
Employee(SSN, E-Name, Office)



What do we do when a **1-to-many** relationship set has an attribute?

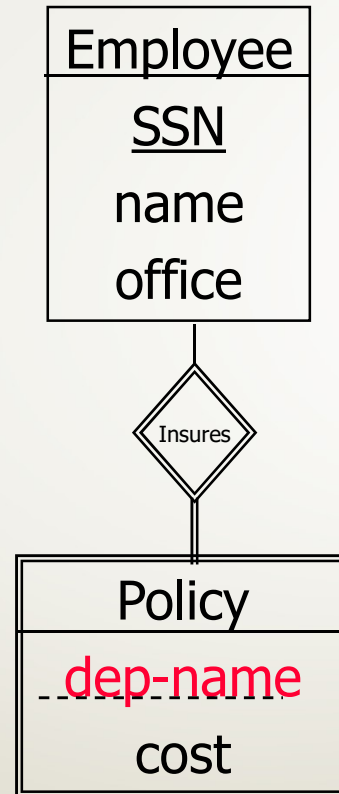


Project(P-number, P-name, Due-Date, MgrSSN,
start-date, end-date)

Employee(SSN, E-Name, Office)



Weak Entity Sets



supporting
entity set

supporting
relationship
set

weak
Entity set



Translating Weak Entity Sets

- Weak entity sets and supporting relationship sets are translated into a single table. Must include **key of supporting entity set**, as a foreign key.
- When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Insurance_Policy (  
  dep-name      CHAR(20),  
  cost          REAL,  
  ssn          CHAR(11) NOT NULL,  
  PRIMARY KEY  (dep-name, ssn),  
  
  FOREIGN KEY (ssn) REFERENCES Employee,  
  ON DELETE CASCADE)
```




Converting Subclass Structures to Relations

Three approaches:

1. Follow the *E /R* viewpoint.

- For each entity set E in the hierarchy, create a relation that includes the key attributes from the root and any attributes belonging to E .

2. Treat entities as objects belonging to a single class.

- For each possible subtree that includes the root, create one relation, whose schema includes all the attributes of all the entity sets in the subtree.

3. Use null values.

- Create one relation with all the attributes of all the entity sets in the hierarchy.
- Each entity is represented by one tuple, and that tuple has a null value for whatever attributes the entity does not have.

Read textbook section 4.6
for examples and comparison of approaches